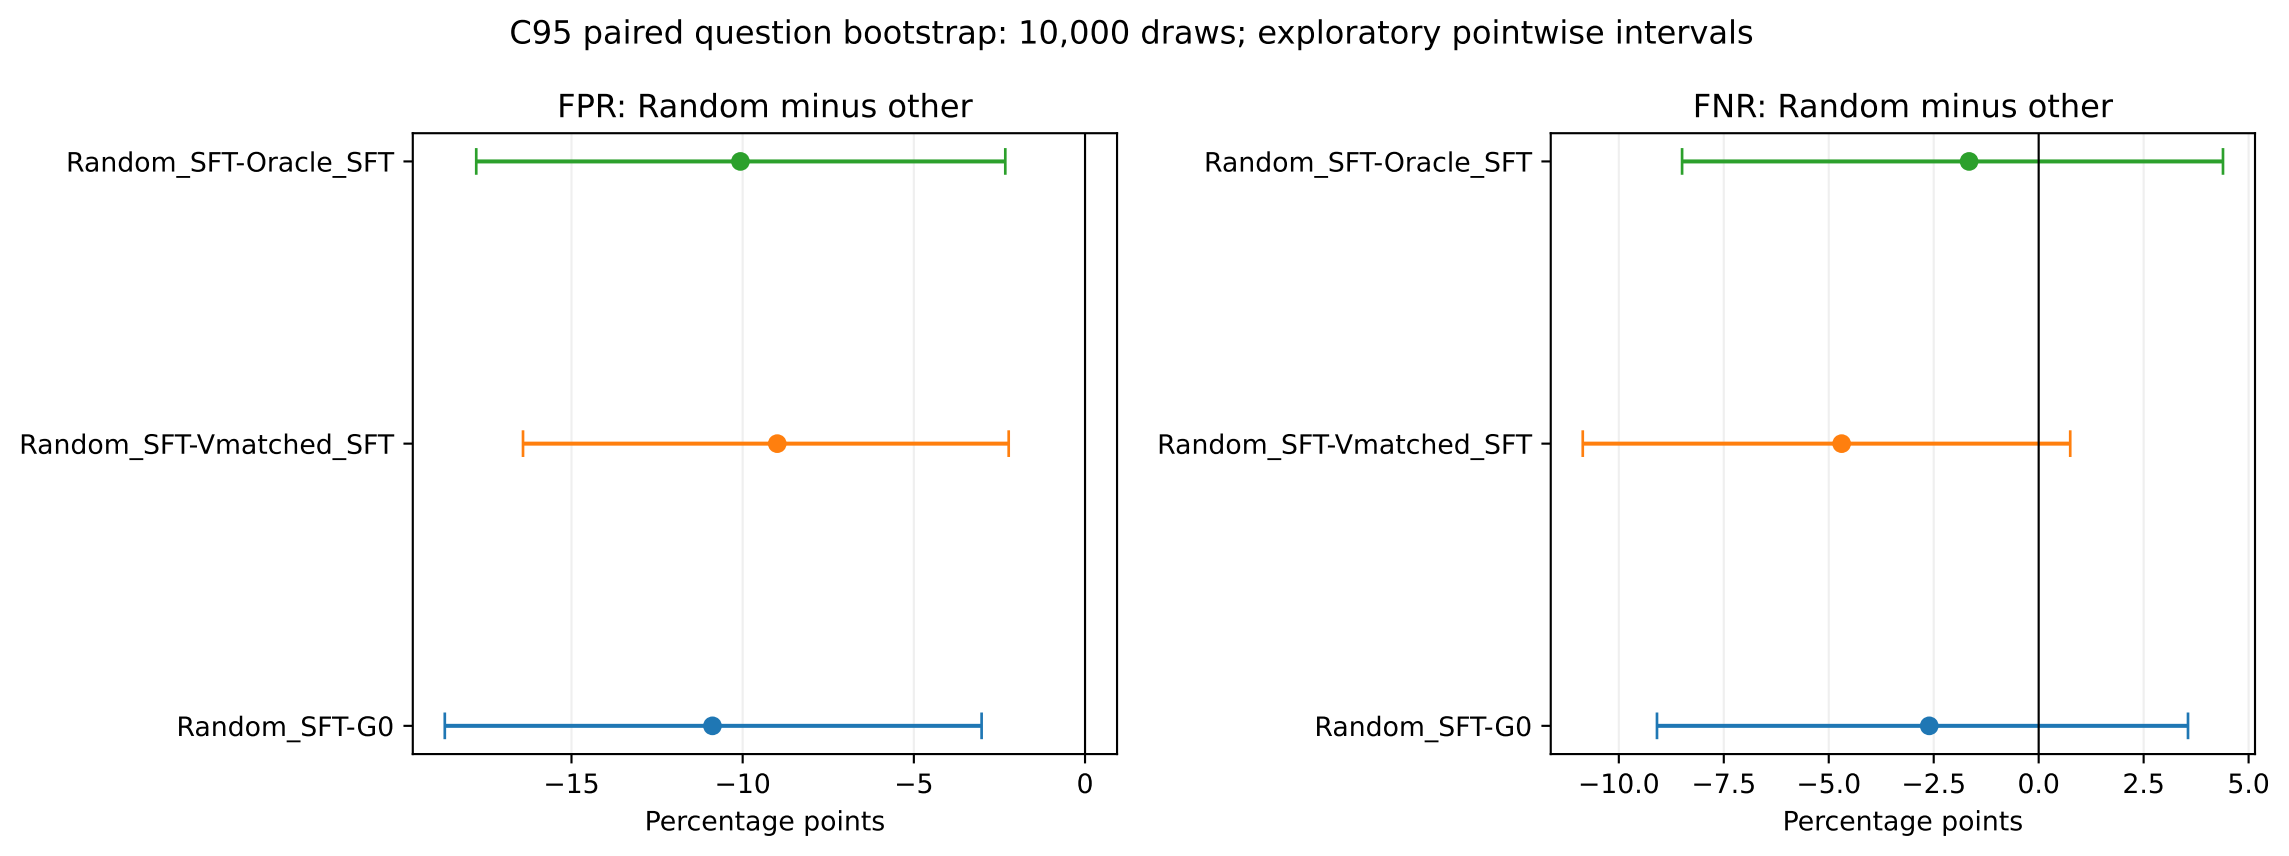
C95 paired question bootstrap: 10,000 draws; exploratory pointwise intervals
FPR: Random minus other
Random_SFT-Oracle_SFT
Random_SFT-Vmatched_SFT
Random_SFT-G0
−15
−10
−5
0
Percentage points
FNR: Random minus other
Random_SFT-Oracle_SFT
Random_SFT-Vmatched_SFT
Random_SFT-G0
−10.0
−7.5
−5.0
−2.5
0.0
2.5
5.0
Percentage points

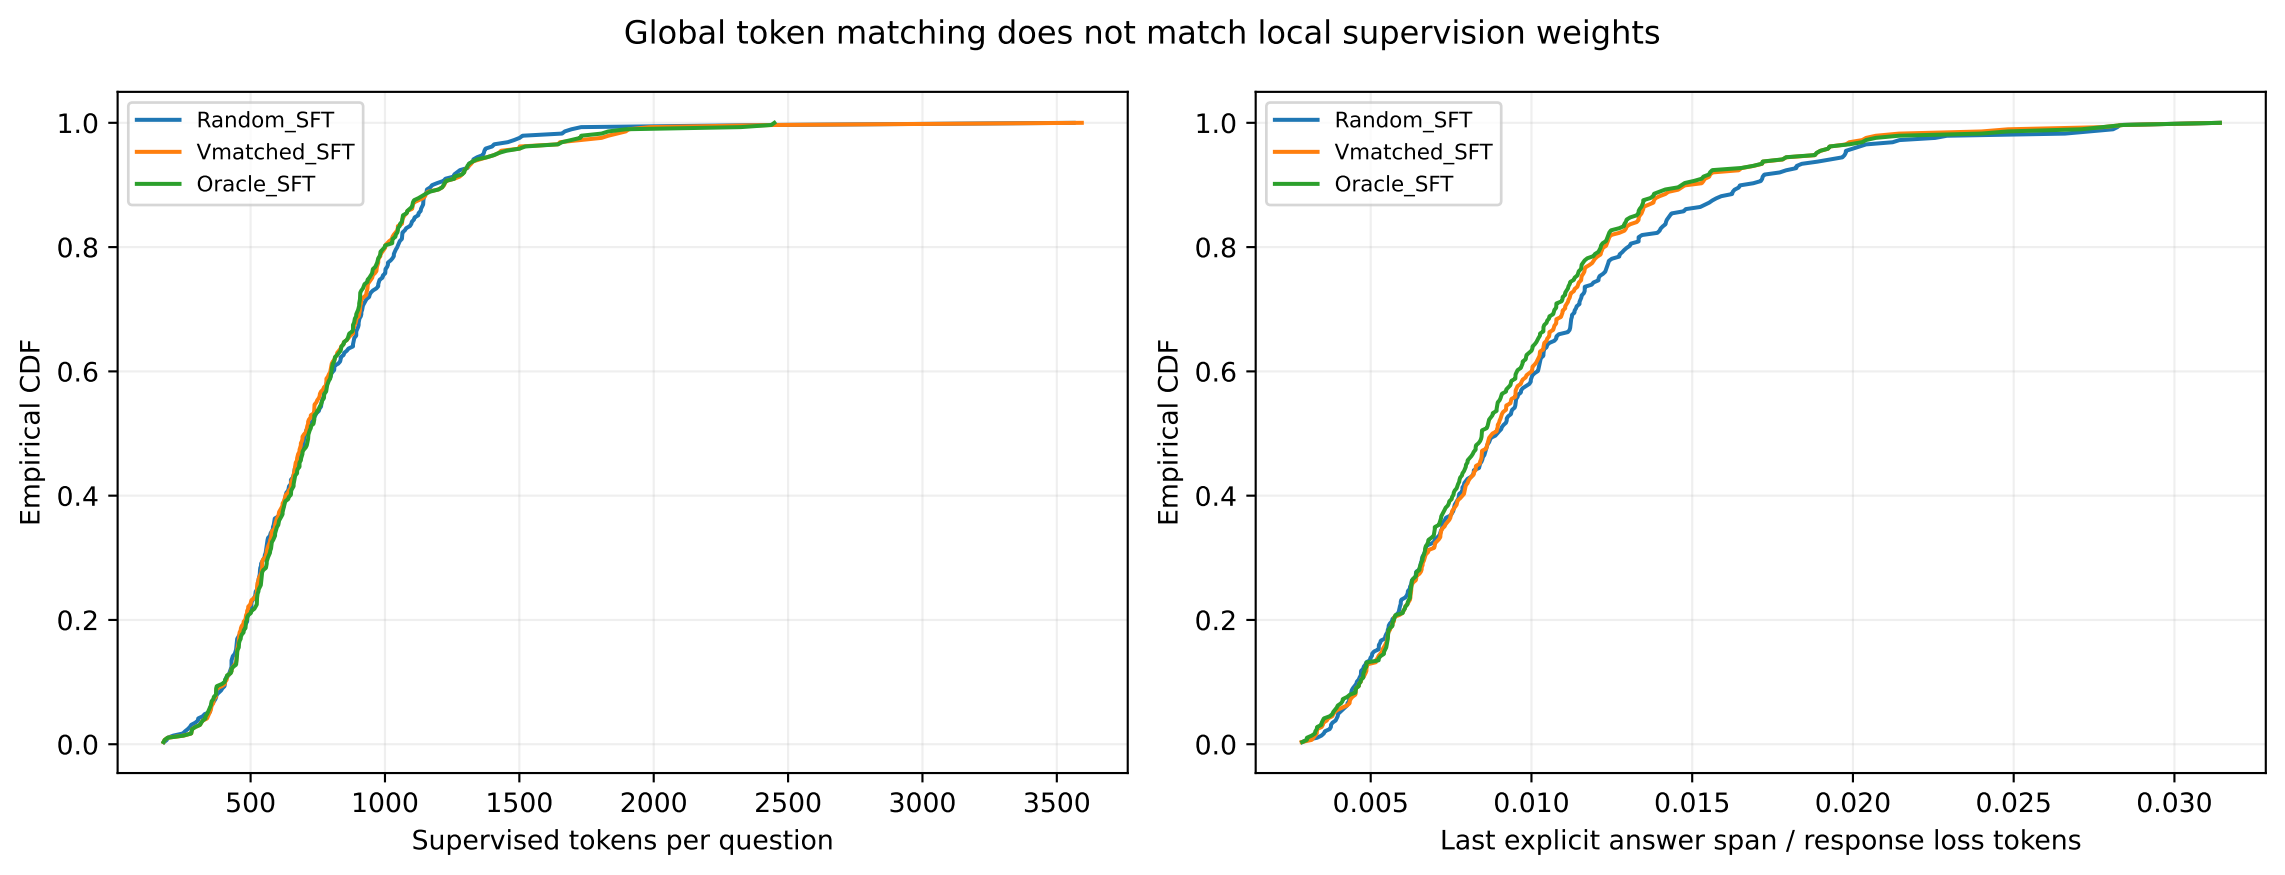
Global token matching does not match local supervision weights
Random_SFT
Vmatched_SFT
Oracle_SFT
Empirical CDF
0.0
0.2
0.4
0.6
0.8
1.0
500
1000
1500
2000
2500
3000
3500
Supervised tokens per question
Random_SFT
Vmatched_SFT
Oracle_SFT
Empirical CDF
0.0
0.2
0.4
0.6
0.8
1.0
0.005
0.010
0.015
0.020
0.025
0.030
Last explicit answer span / response loss tokens

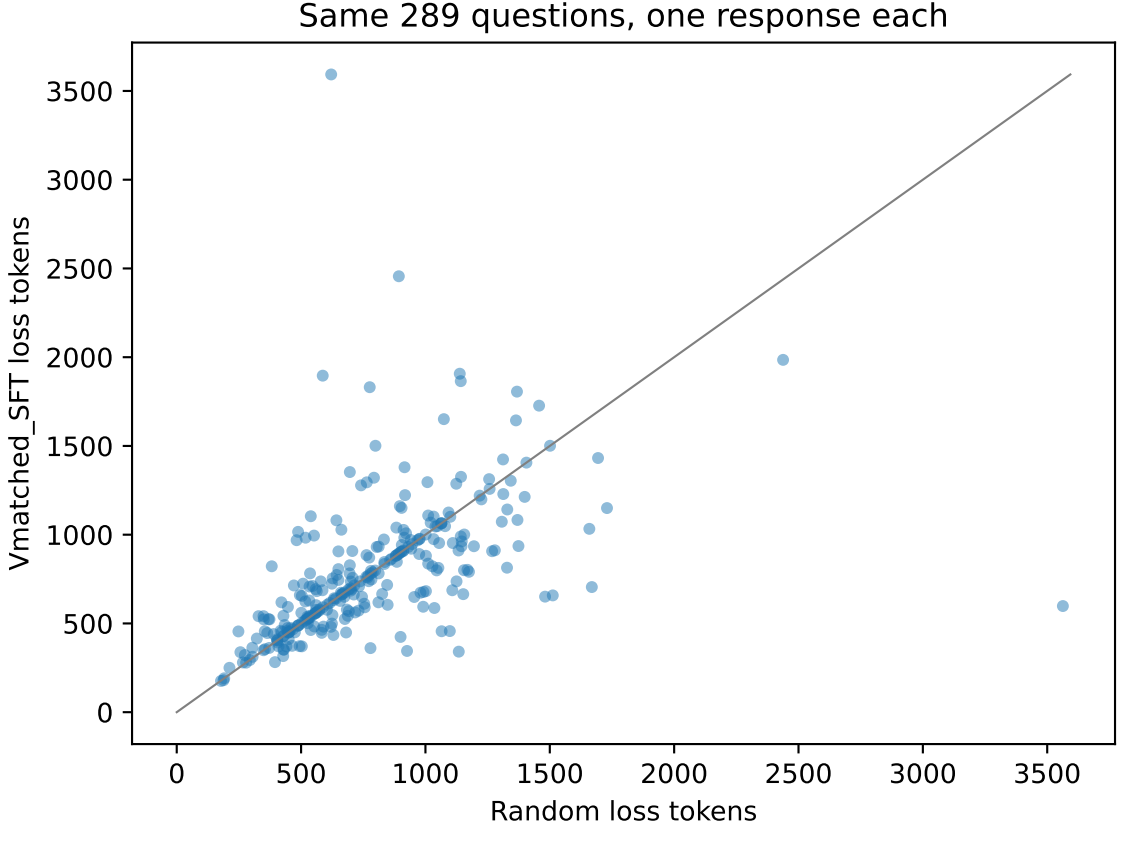
Same 289 questions, one response each
Vmatched_SFT loss tokens
Random loss tokens
0
500
1000
1500
2000
2500
3000
3500

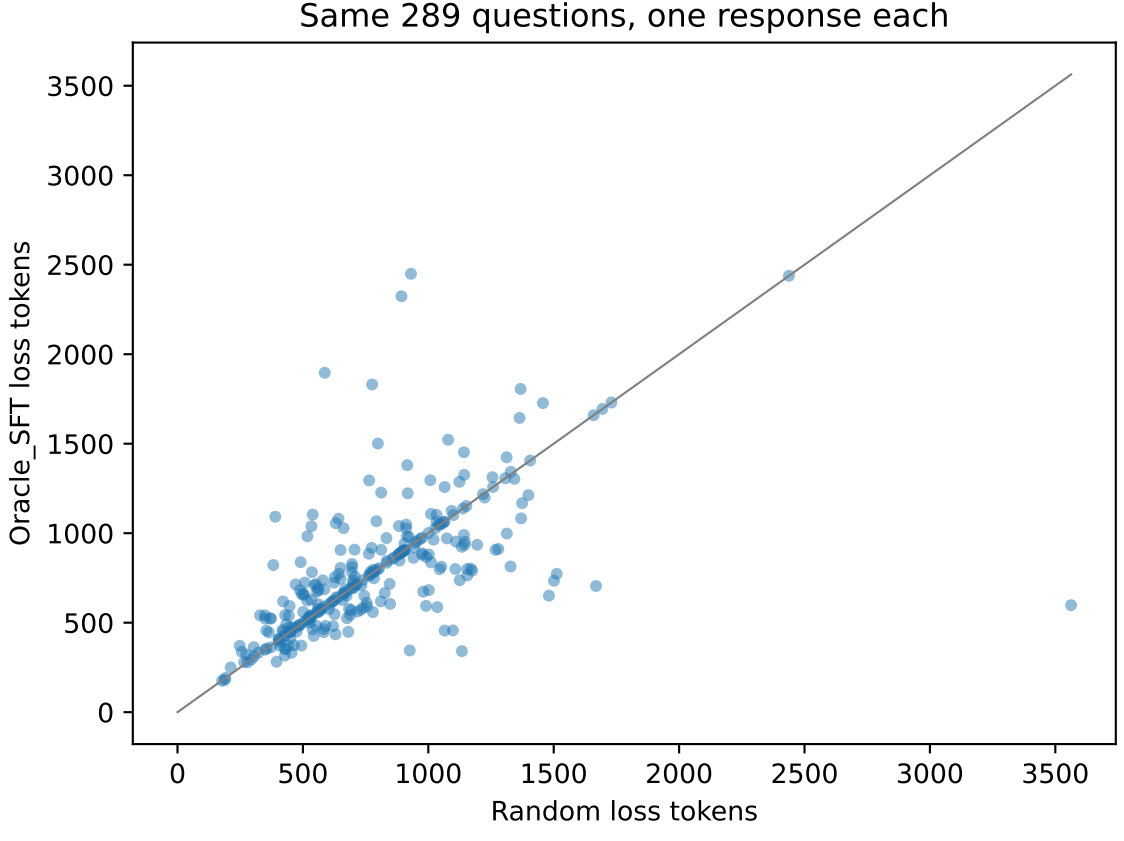
Same 289 questions, one response each
Oracle_SFT loss tokens
Random loss tokens
0
500
1000
1500
2000
2500
3000
3500