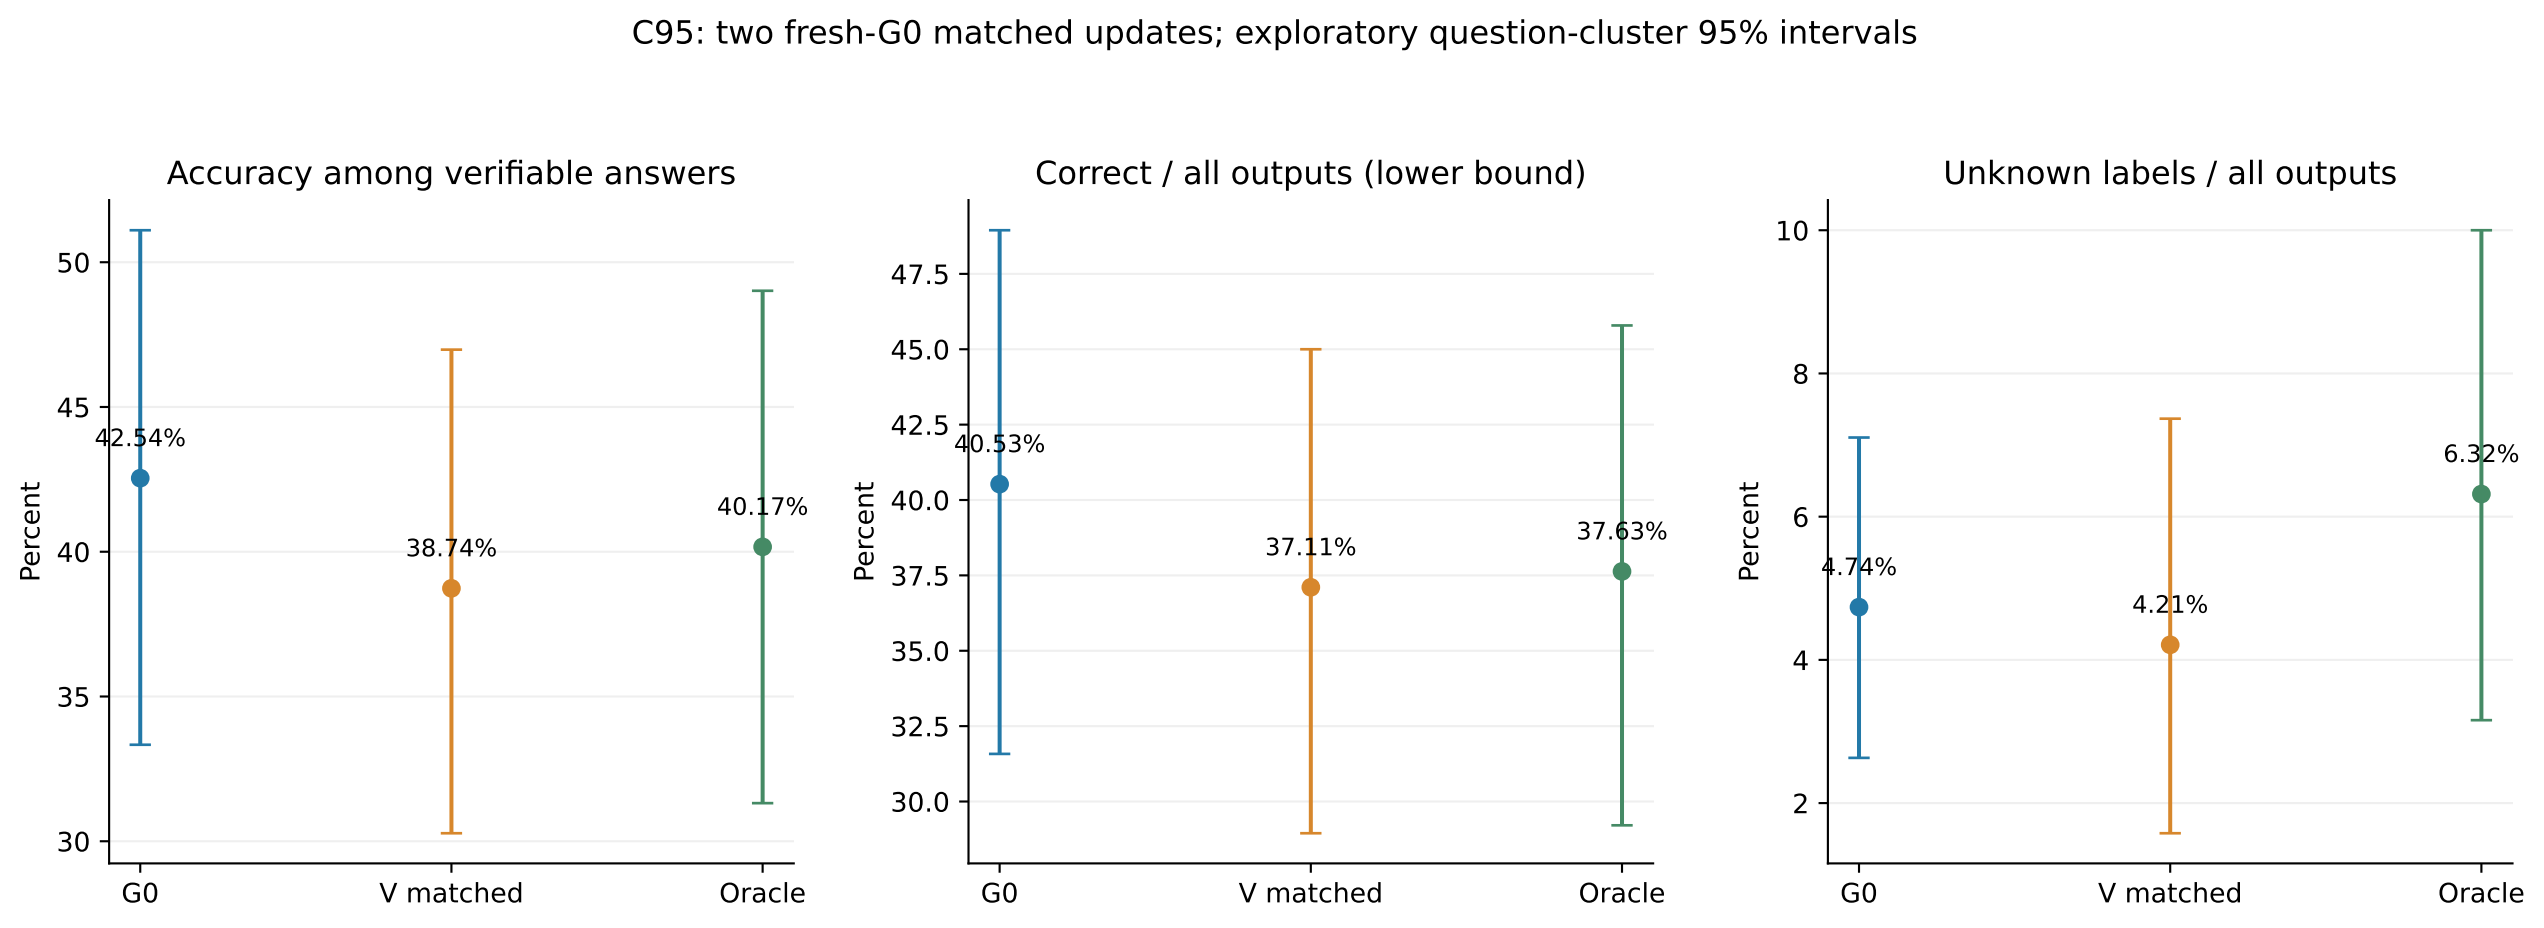
C95: two fresh-G0 matched updates; exploratory question-cluster 95% intervals
Accuracy among verifiable answers
Correct / all outputs (lower bound)
Unknown labels / all outputs
Percent
50
45
40
35
30
47.5
45.0
42.5
40.0
37.5
35.0
32.5
30.0
10
8
6
4
2
42.54%
38.74%
40.17%
40.53%
37.11%
37.63%
4.74%
4.21%
6.32%
G0
V matched
Oracle

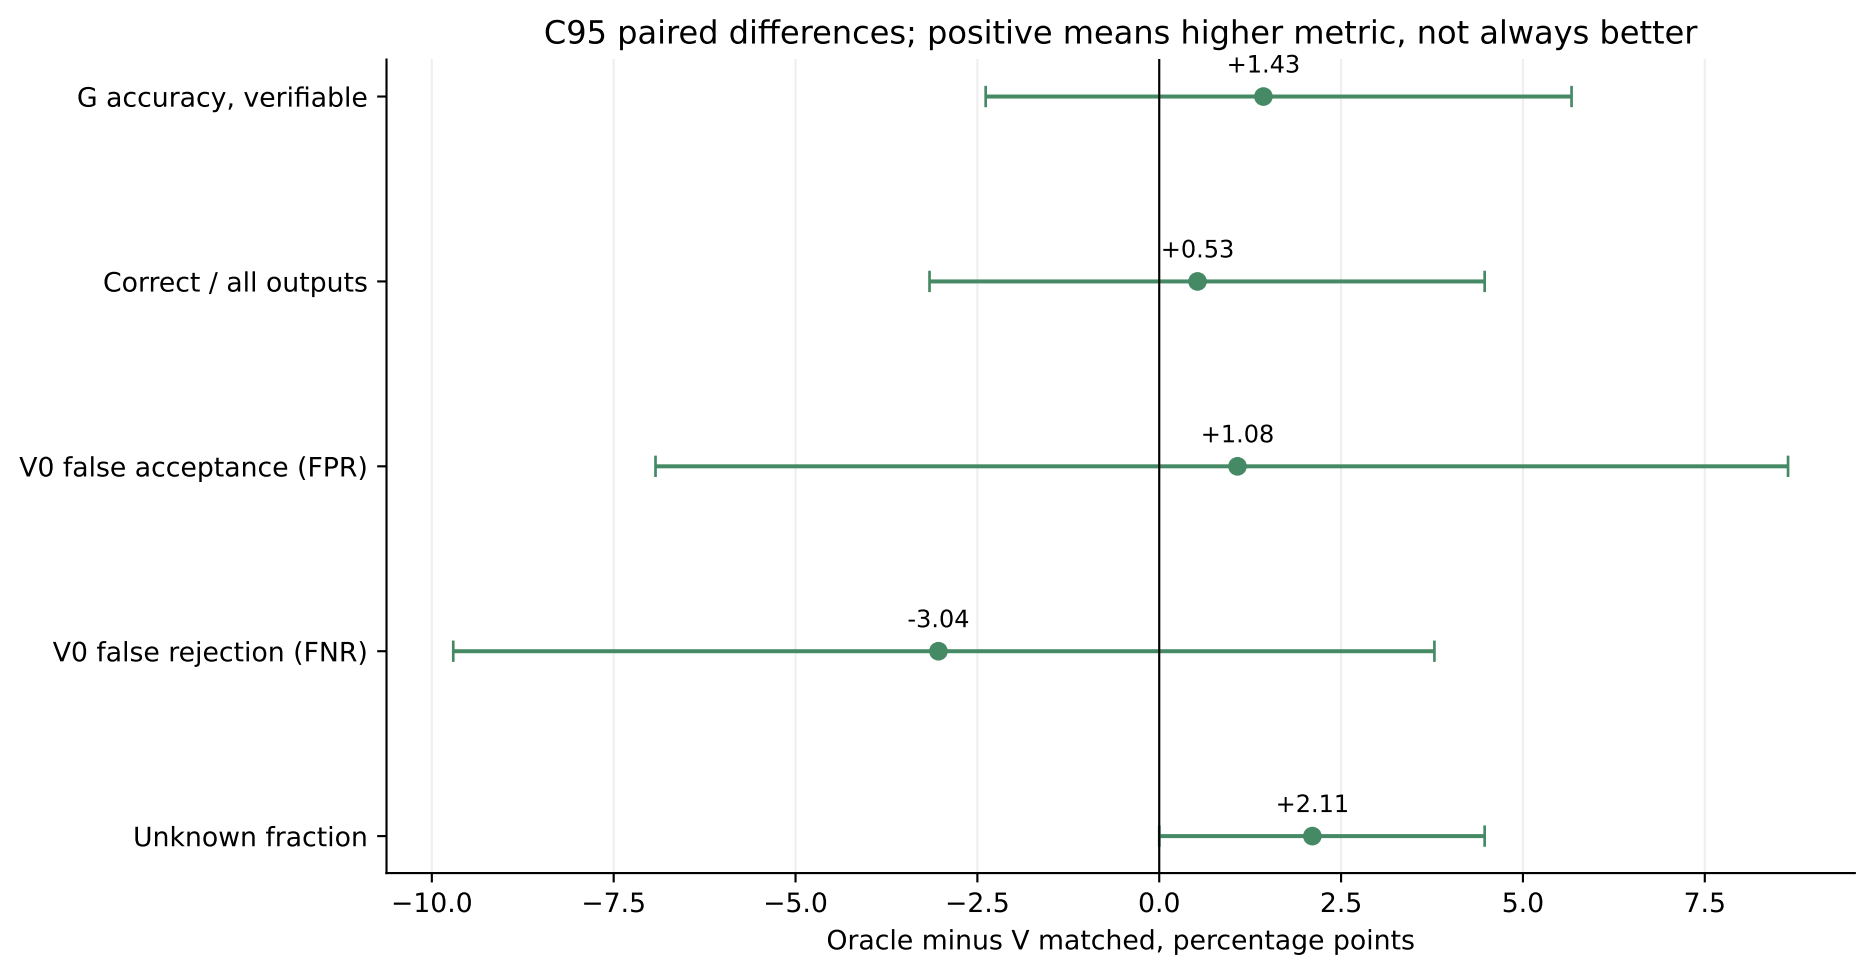
C95 paired differences; positive means higher metric, not always better
G accuracy, verifiable
Correct / all outputs
V0 false acceptance (FPR)
V0 false rejection (FNR)
Unknown fraction
+1.43
+0.53
+1.08
-3.04
+2.11
−10.0
−7.5
−5.0
−2.5
0.0
2.5
5.0
7.5
Oracle minus V matched, percentage points

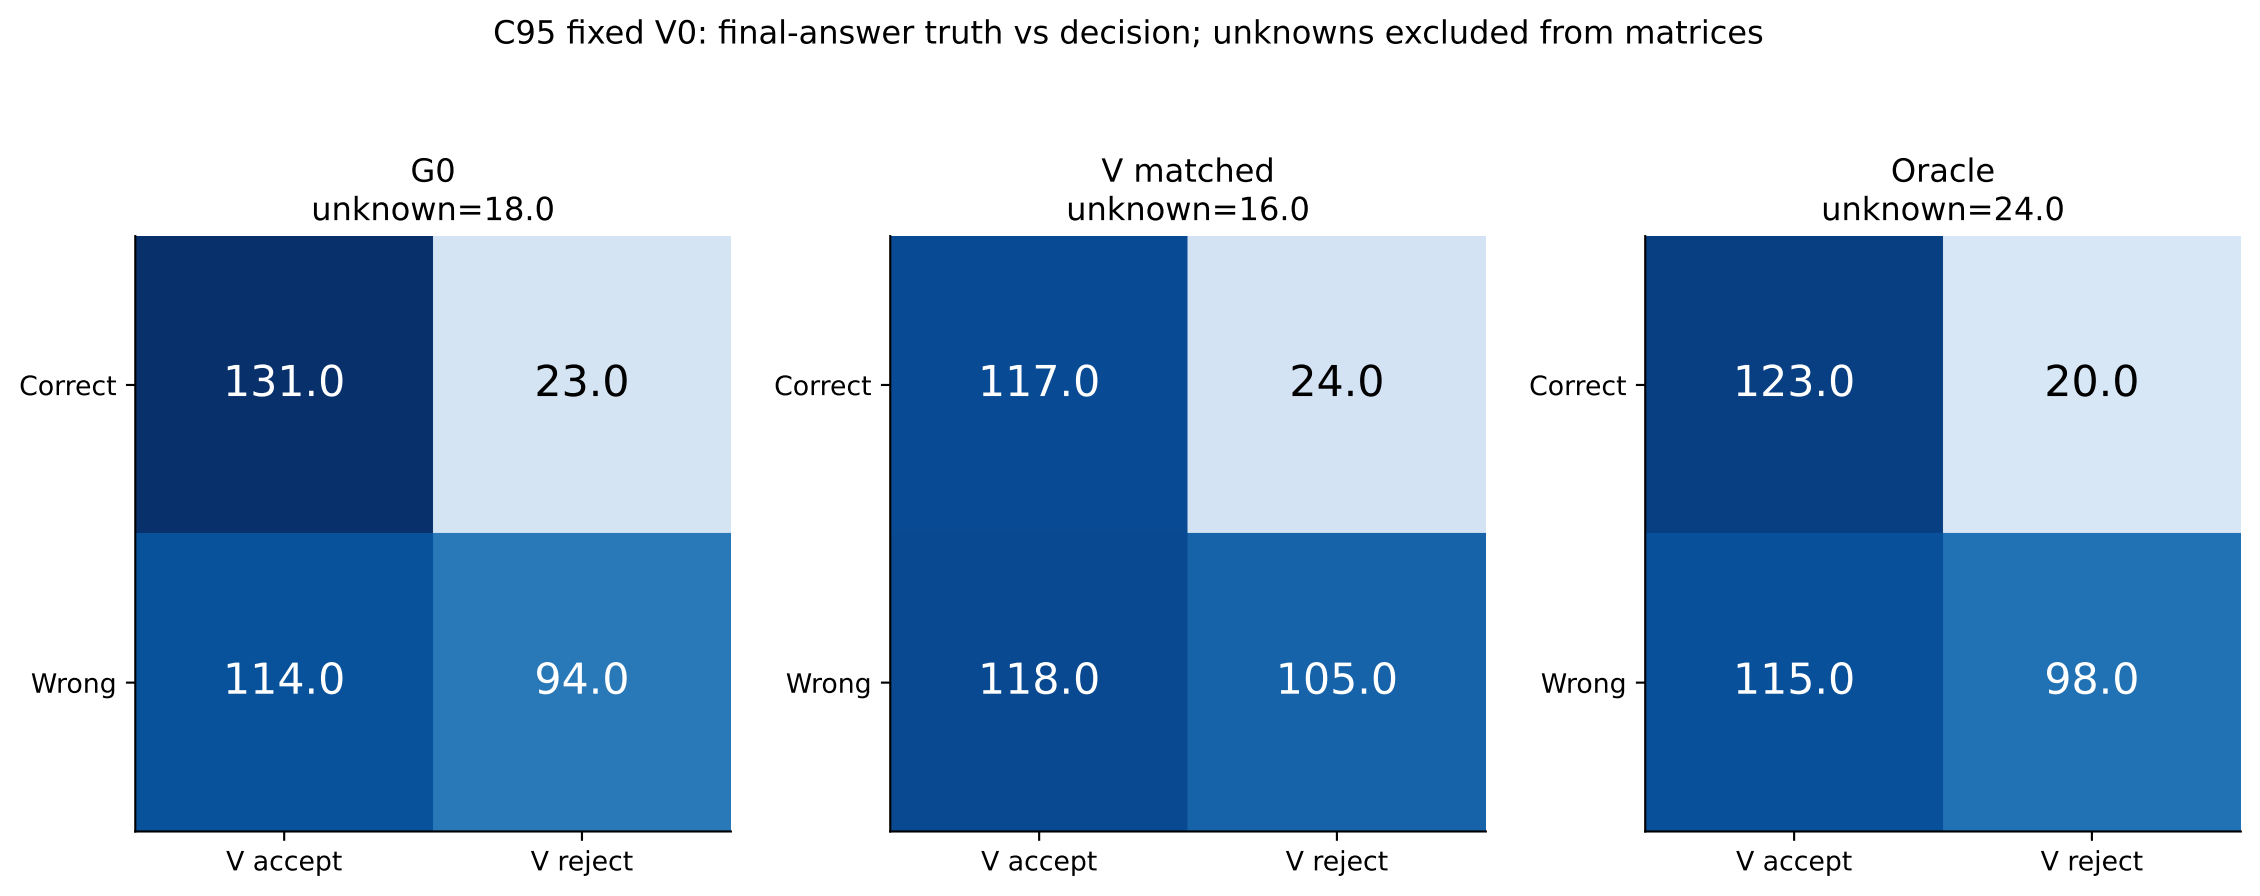
C95 fixed V0: final-answer truth vs decision; unknowns excluded from matrices
G0
unknown=18.0
Correct
Wrong
131.0
23.0
114.0
94.0
V accept
V reject
V matched
unknown=16.0
Correct
Wrong
117.0
24.0
118.0
105.0
V accept
V reject
Oracle
unknown=24.0
Correct
Wrong
123.0
20.0
115.0
98.0
V accept
V reject

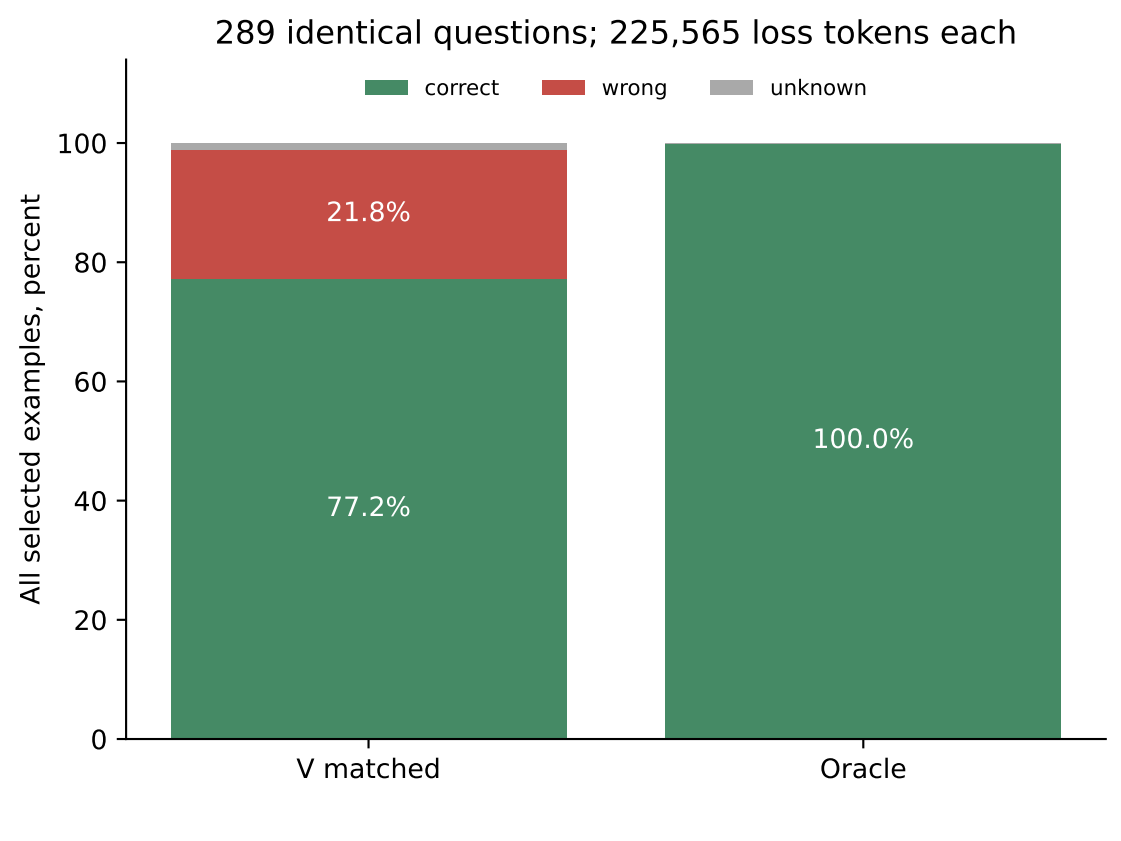
289 identical questions; 225,565 loss tokens each
correct
wrong
unknown
All selected examples, percent
100
80
60
40
20
0
21.8%
77.2%
100.0%
V matched
Oracle

Same 37 steps; training loss is not test accuracy
V matched
Oracle
Mean response-only SFT loss
0.08
0.07
0.06
0.05
0.04
0.03
0
5
10
15
20
25
30
35
Optimizer step

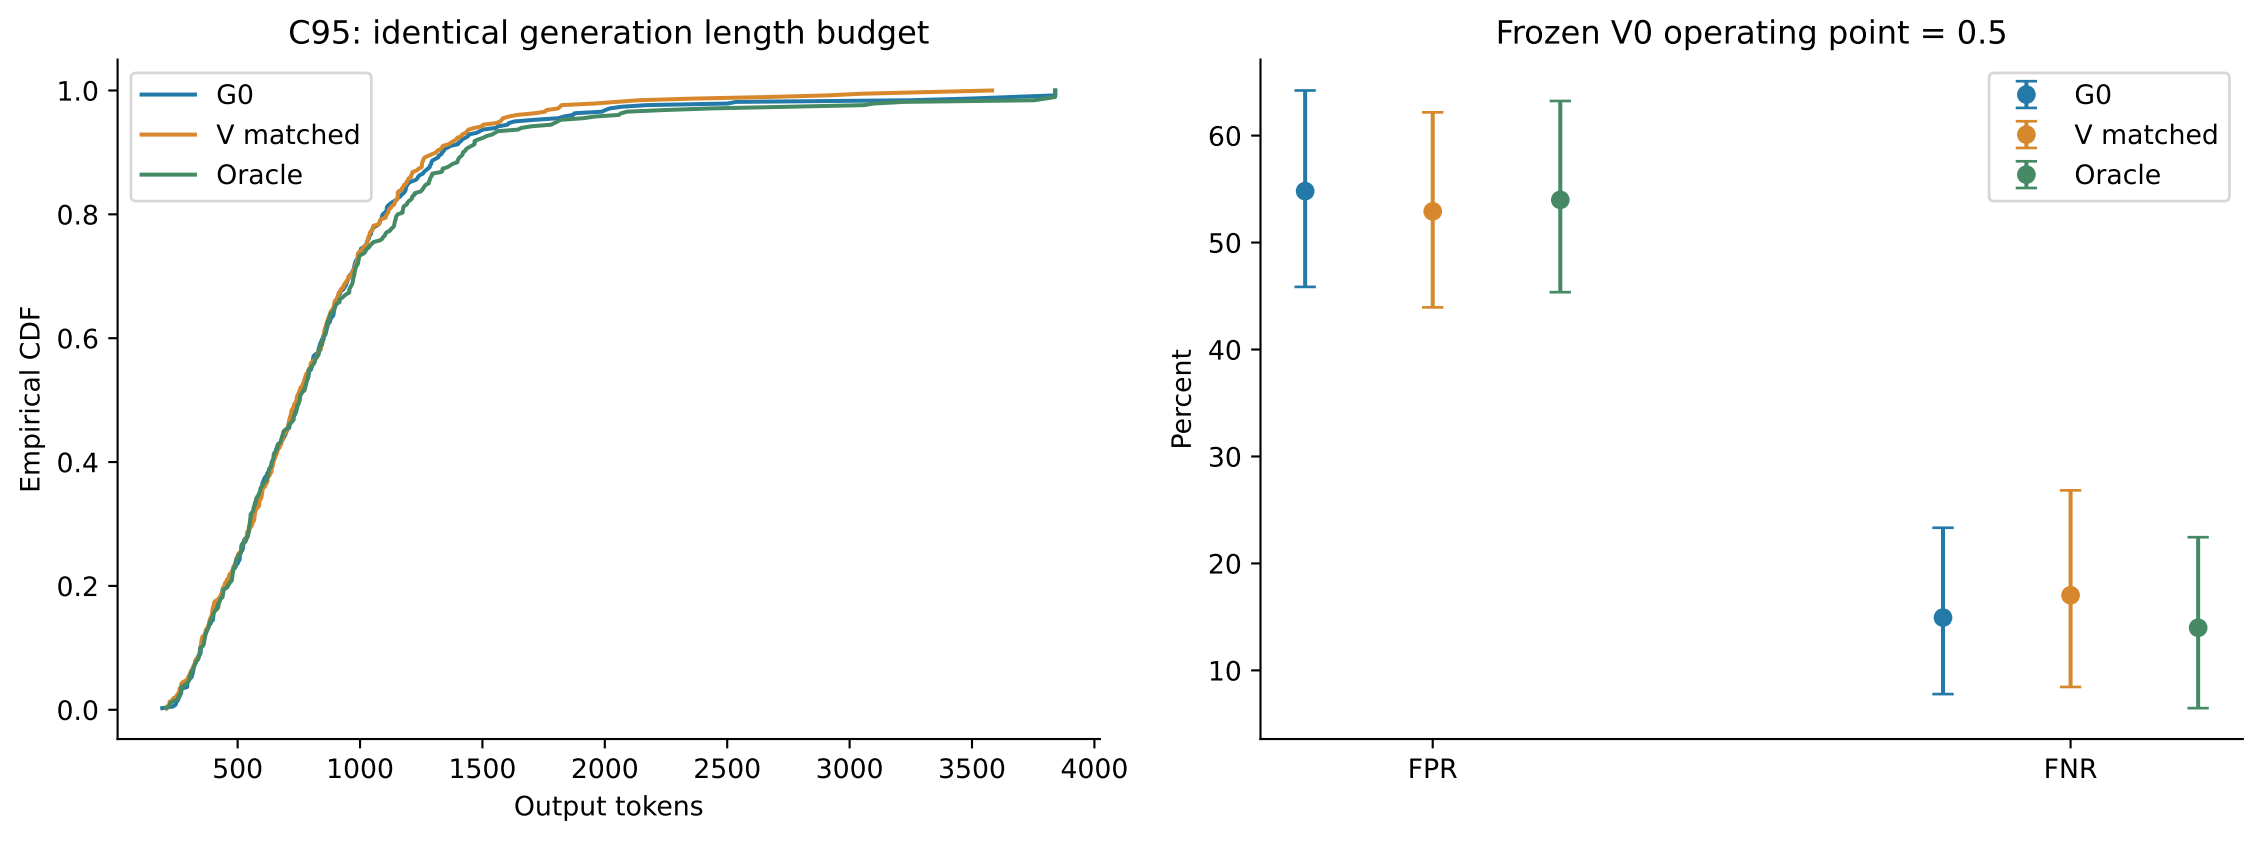
C95: identical generation length budget
G0
V matched
Oracle
Empirical CDF
Output tokens
0.0
0.2
0.4
0.6
0.8
1.0
500
1000
1500
2000
2500
3000
3500
4000
Frozen V0 operating point = 0.5
G0
V matched
Oracle
Percent
10
20
30
40
50
60
FPR
FNR