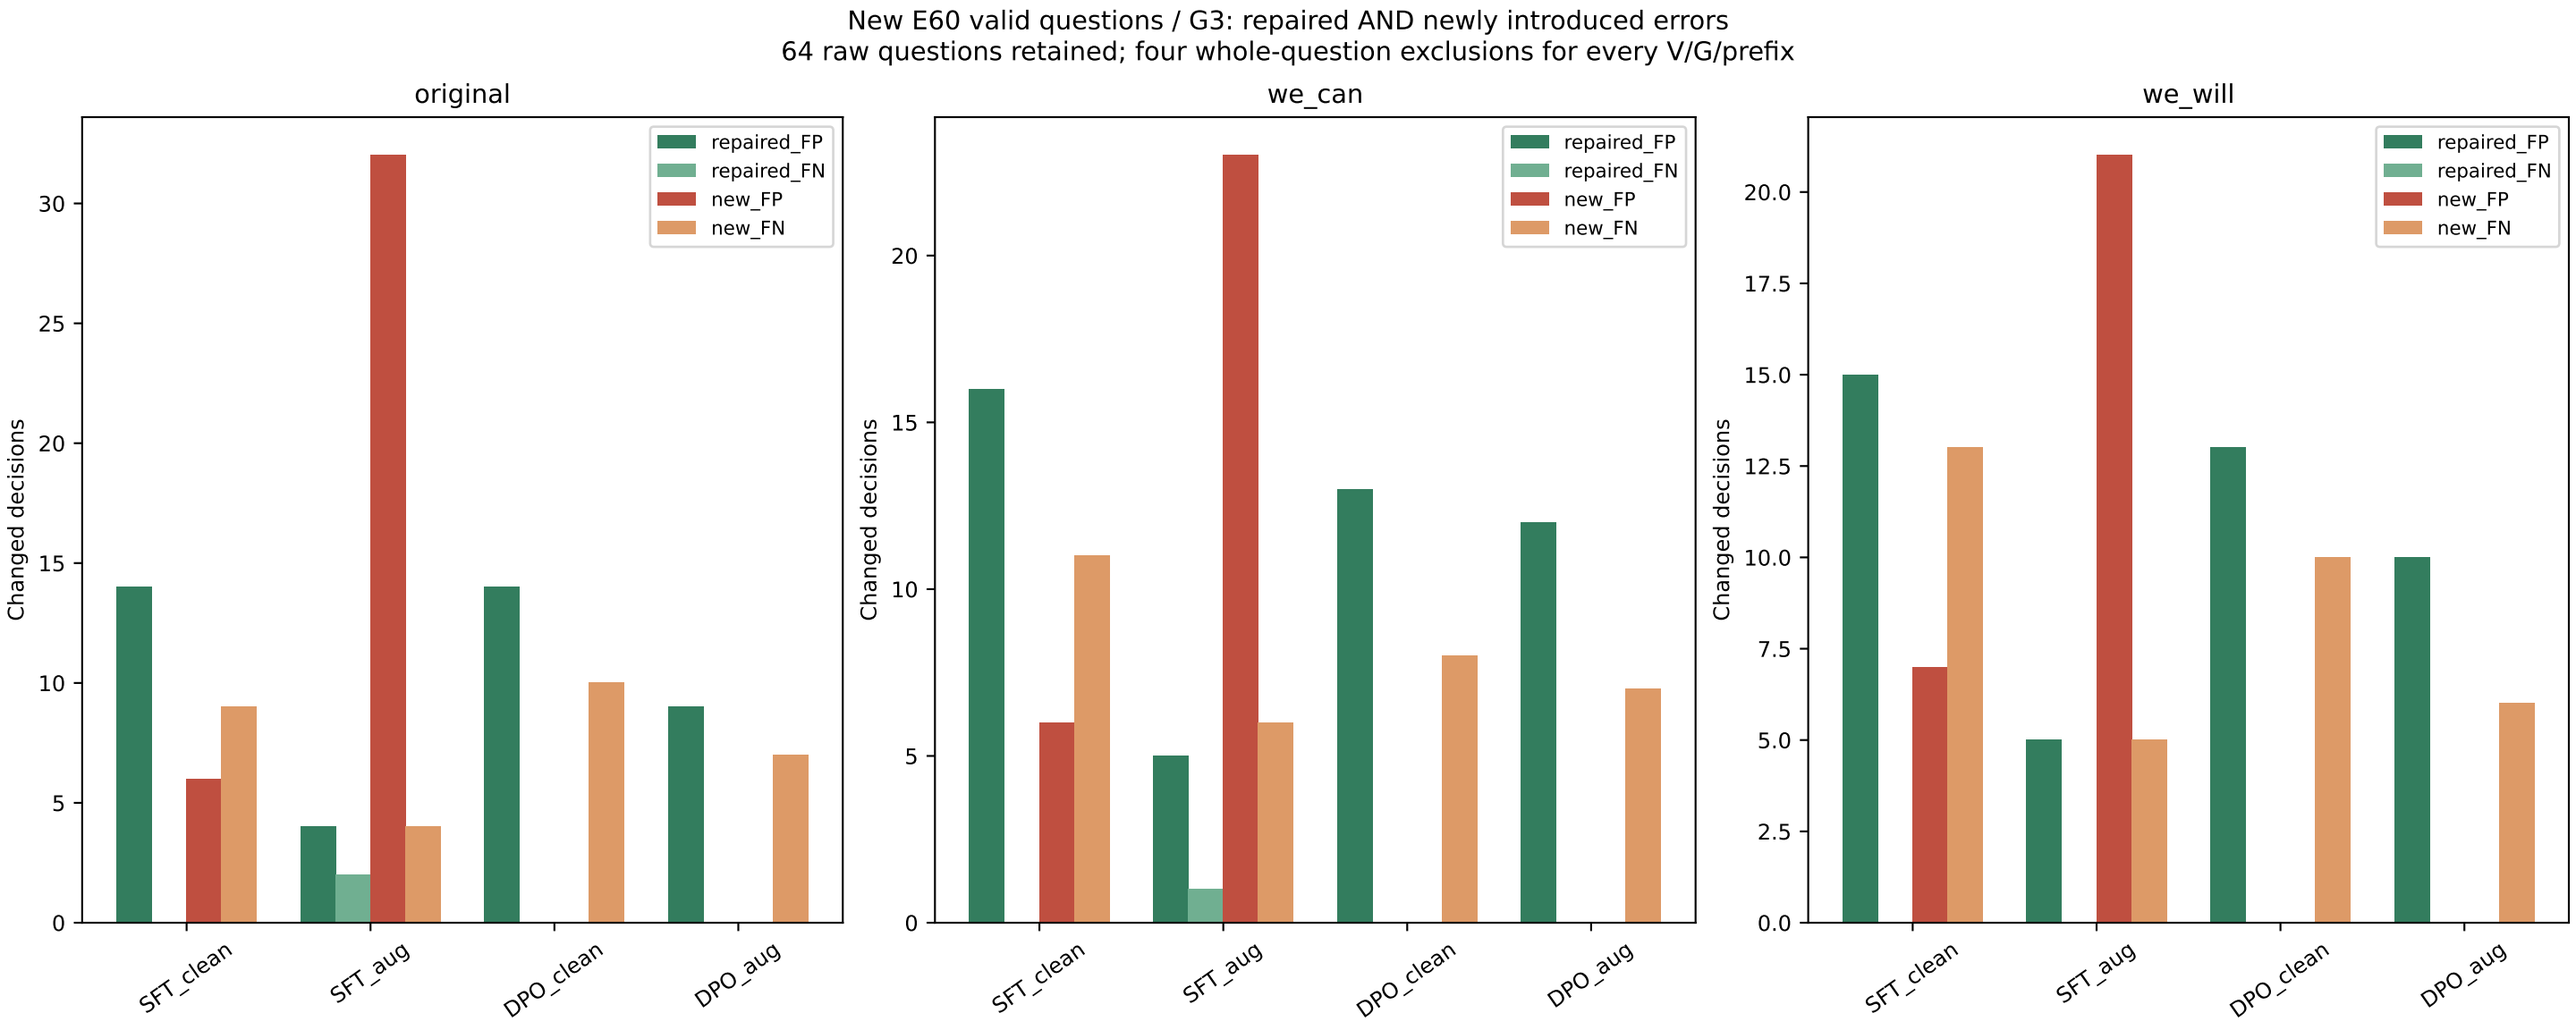
New E60 valid questions / G3: repaired AND newly introduced errors
64 raw questions retained; four whole-question exclusions for every V/G/prefix
original
we_can
we_will
repaired_FP
repaired_FN
new_FP
new_FN
Changed decisions
SFT_clean
SFT_aug
DPO_clean
DPO_aug

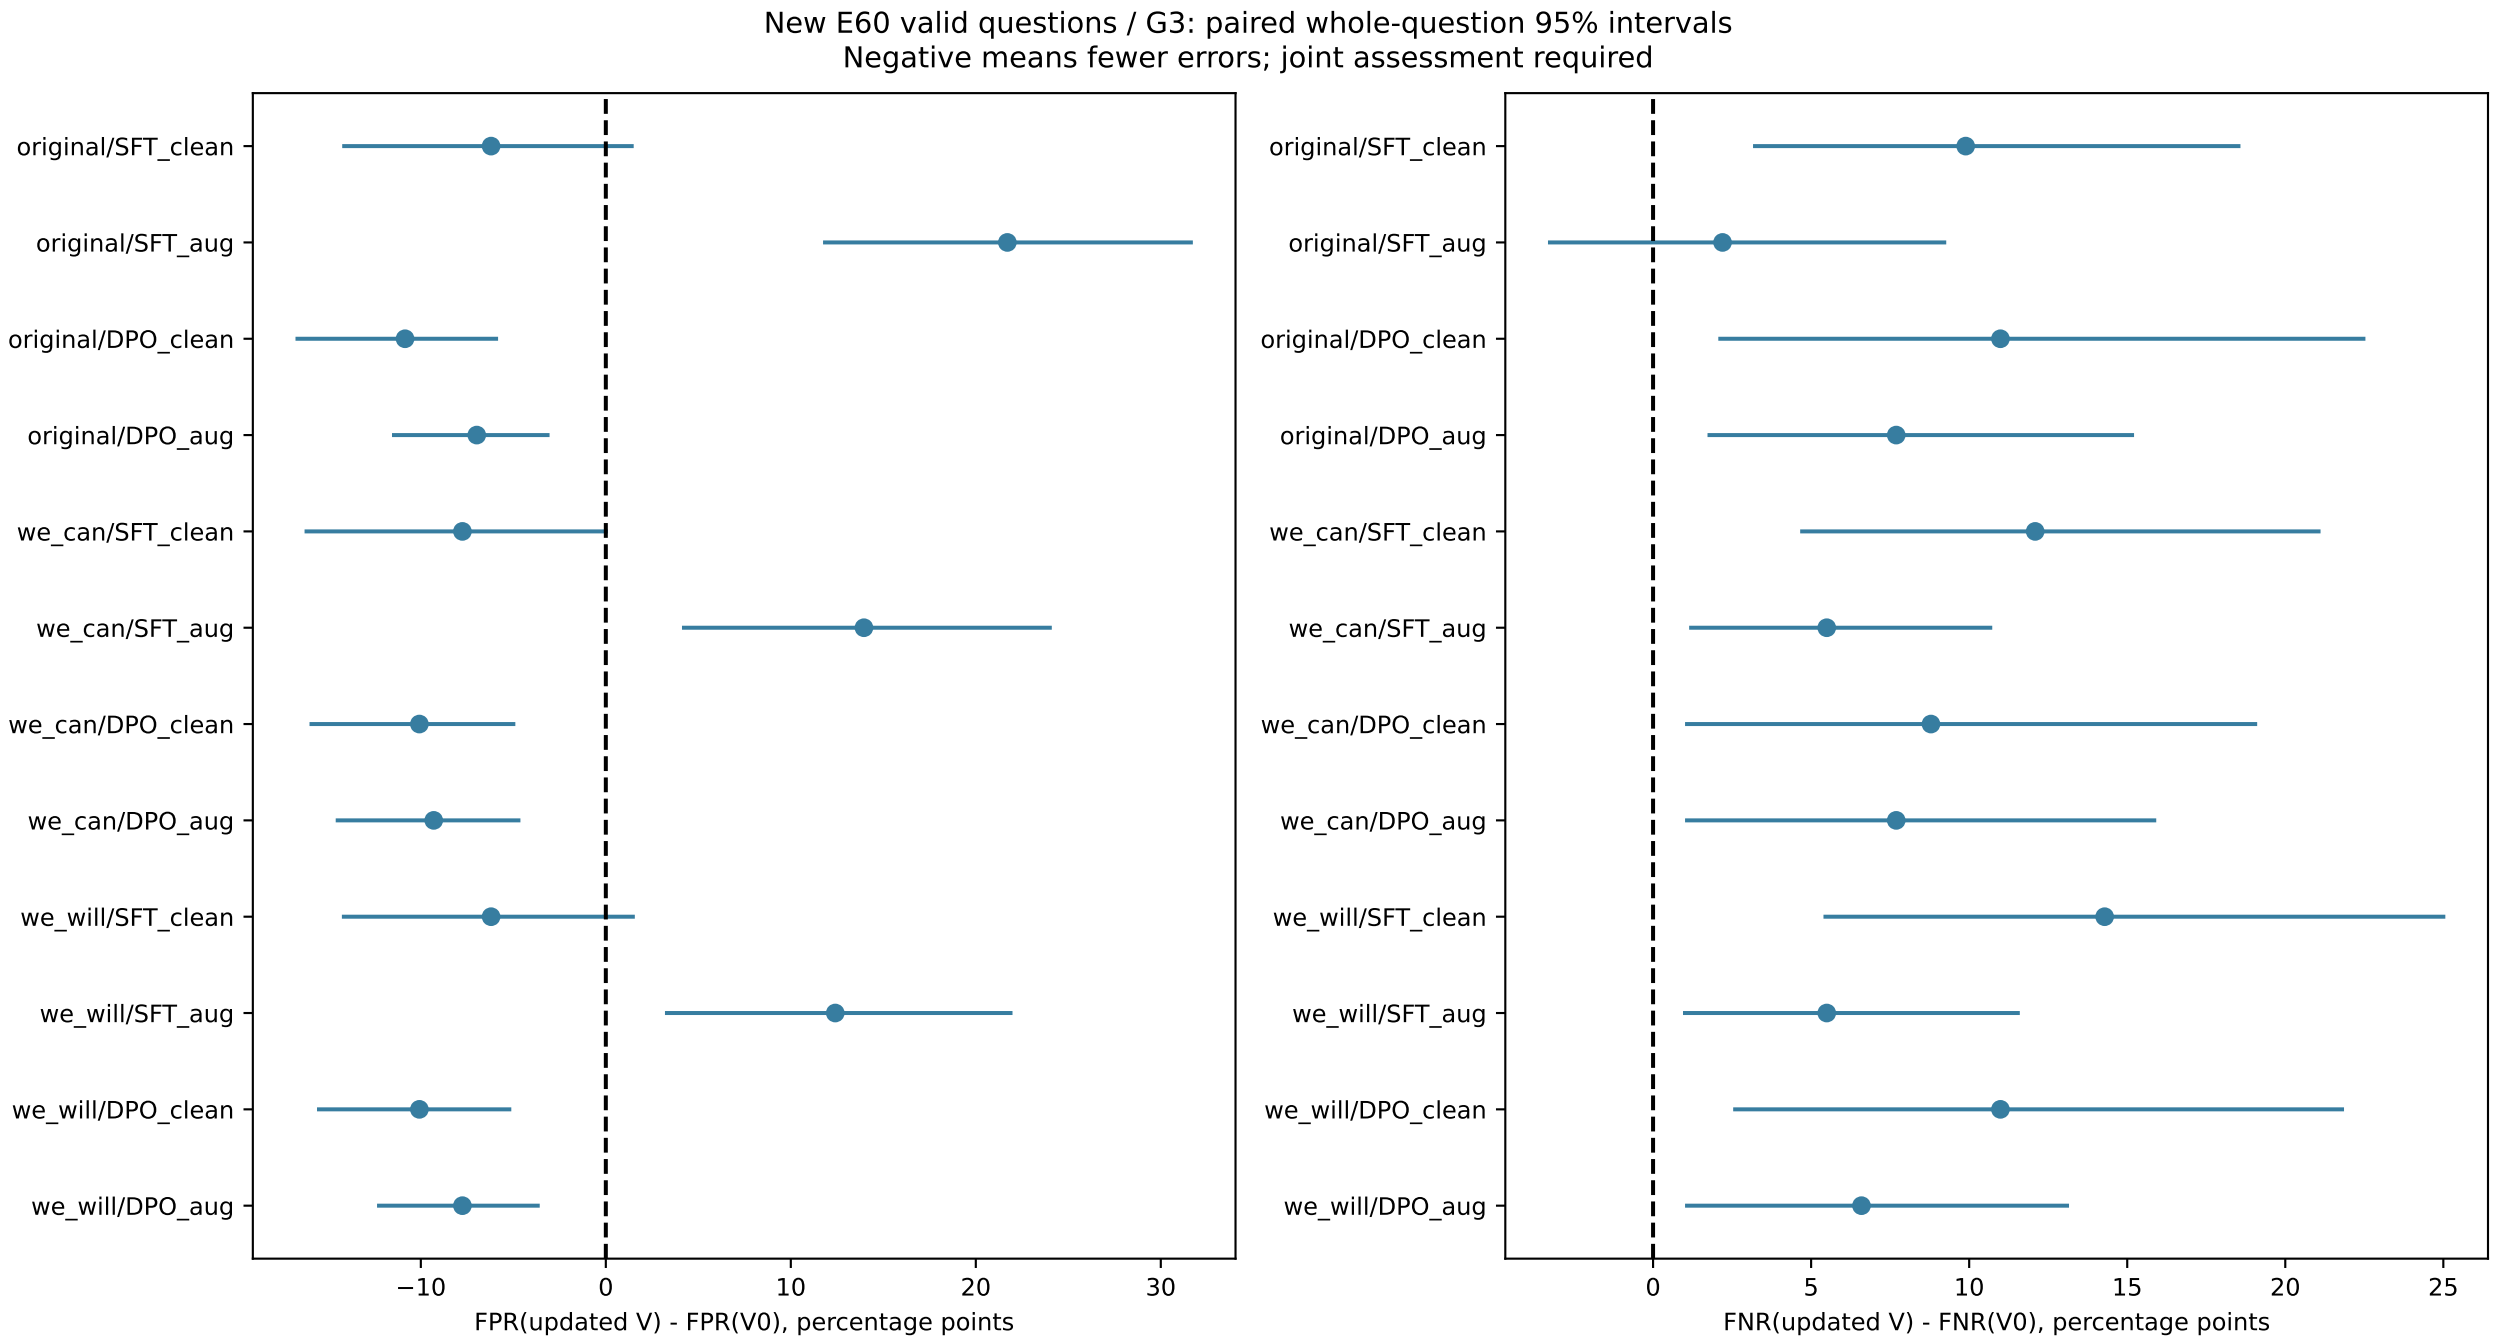
New E60 valid questions / G3: paired whole-question 95% intervals
Negative means fewer errors; joint assessment required
original/SFT_clean
original/SFT_aug
original/DPO_clean
original/DPO_aug
we_can/SFT_clean
we_can/SFT_aug
we_can/DPO_clean
we_can/DPO_aug
we_will/SFT_clean
we_will/SFT_aug
we_will/DPO_clean
we_will/DPO_aug
−10
0
10
20
30
FPR(updated V) - FPR(V0), percentage points
original/SFT_clean
original/SFT_aug
original/DPO_clean
original/DPO_aug
we_can/SFT_clean
we_can/SFT_aug
we_can/DPO_clean
we_can/DPO_aug
we_will/SFT_clean
we_will/SFT_aug
we_will/DPO_clean
we_will/DPO_aug
0
5
10
15
20
25
FNR(updated V) - FNR(V0), percentage points

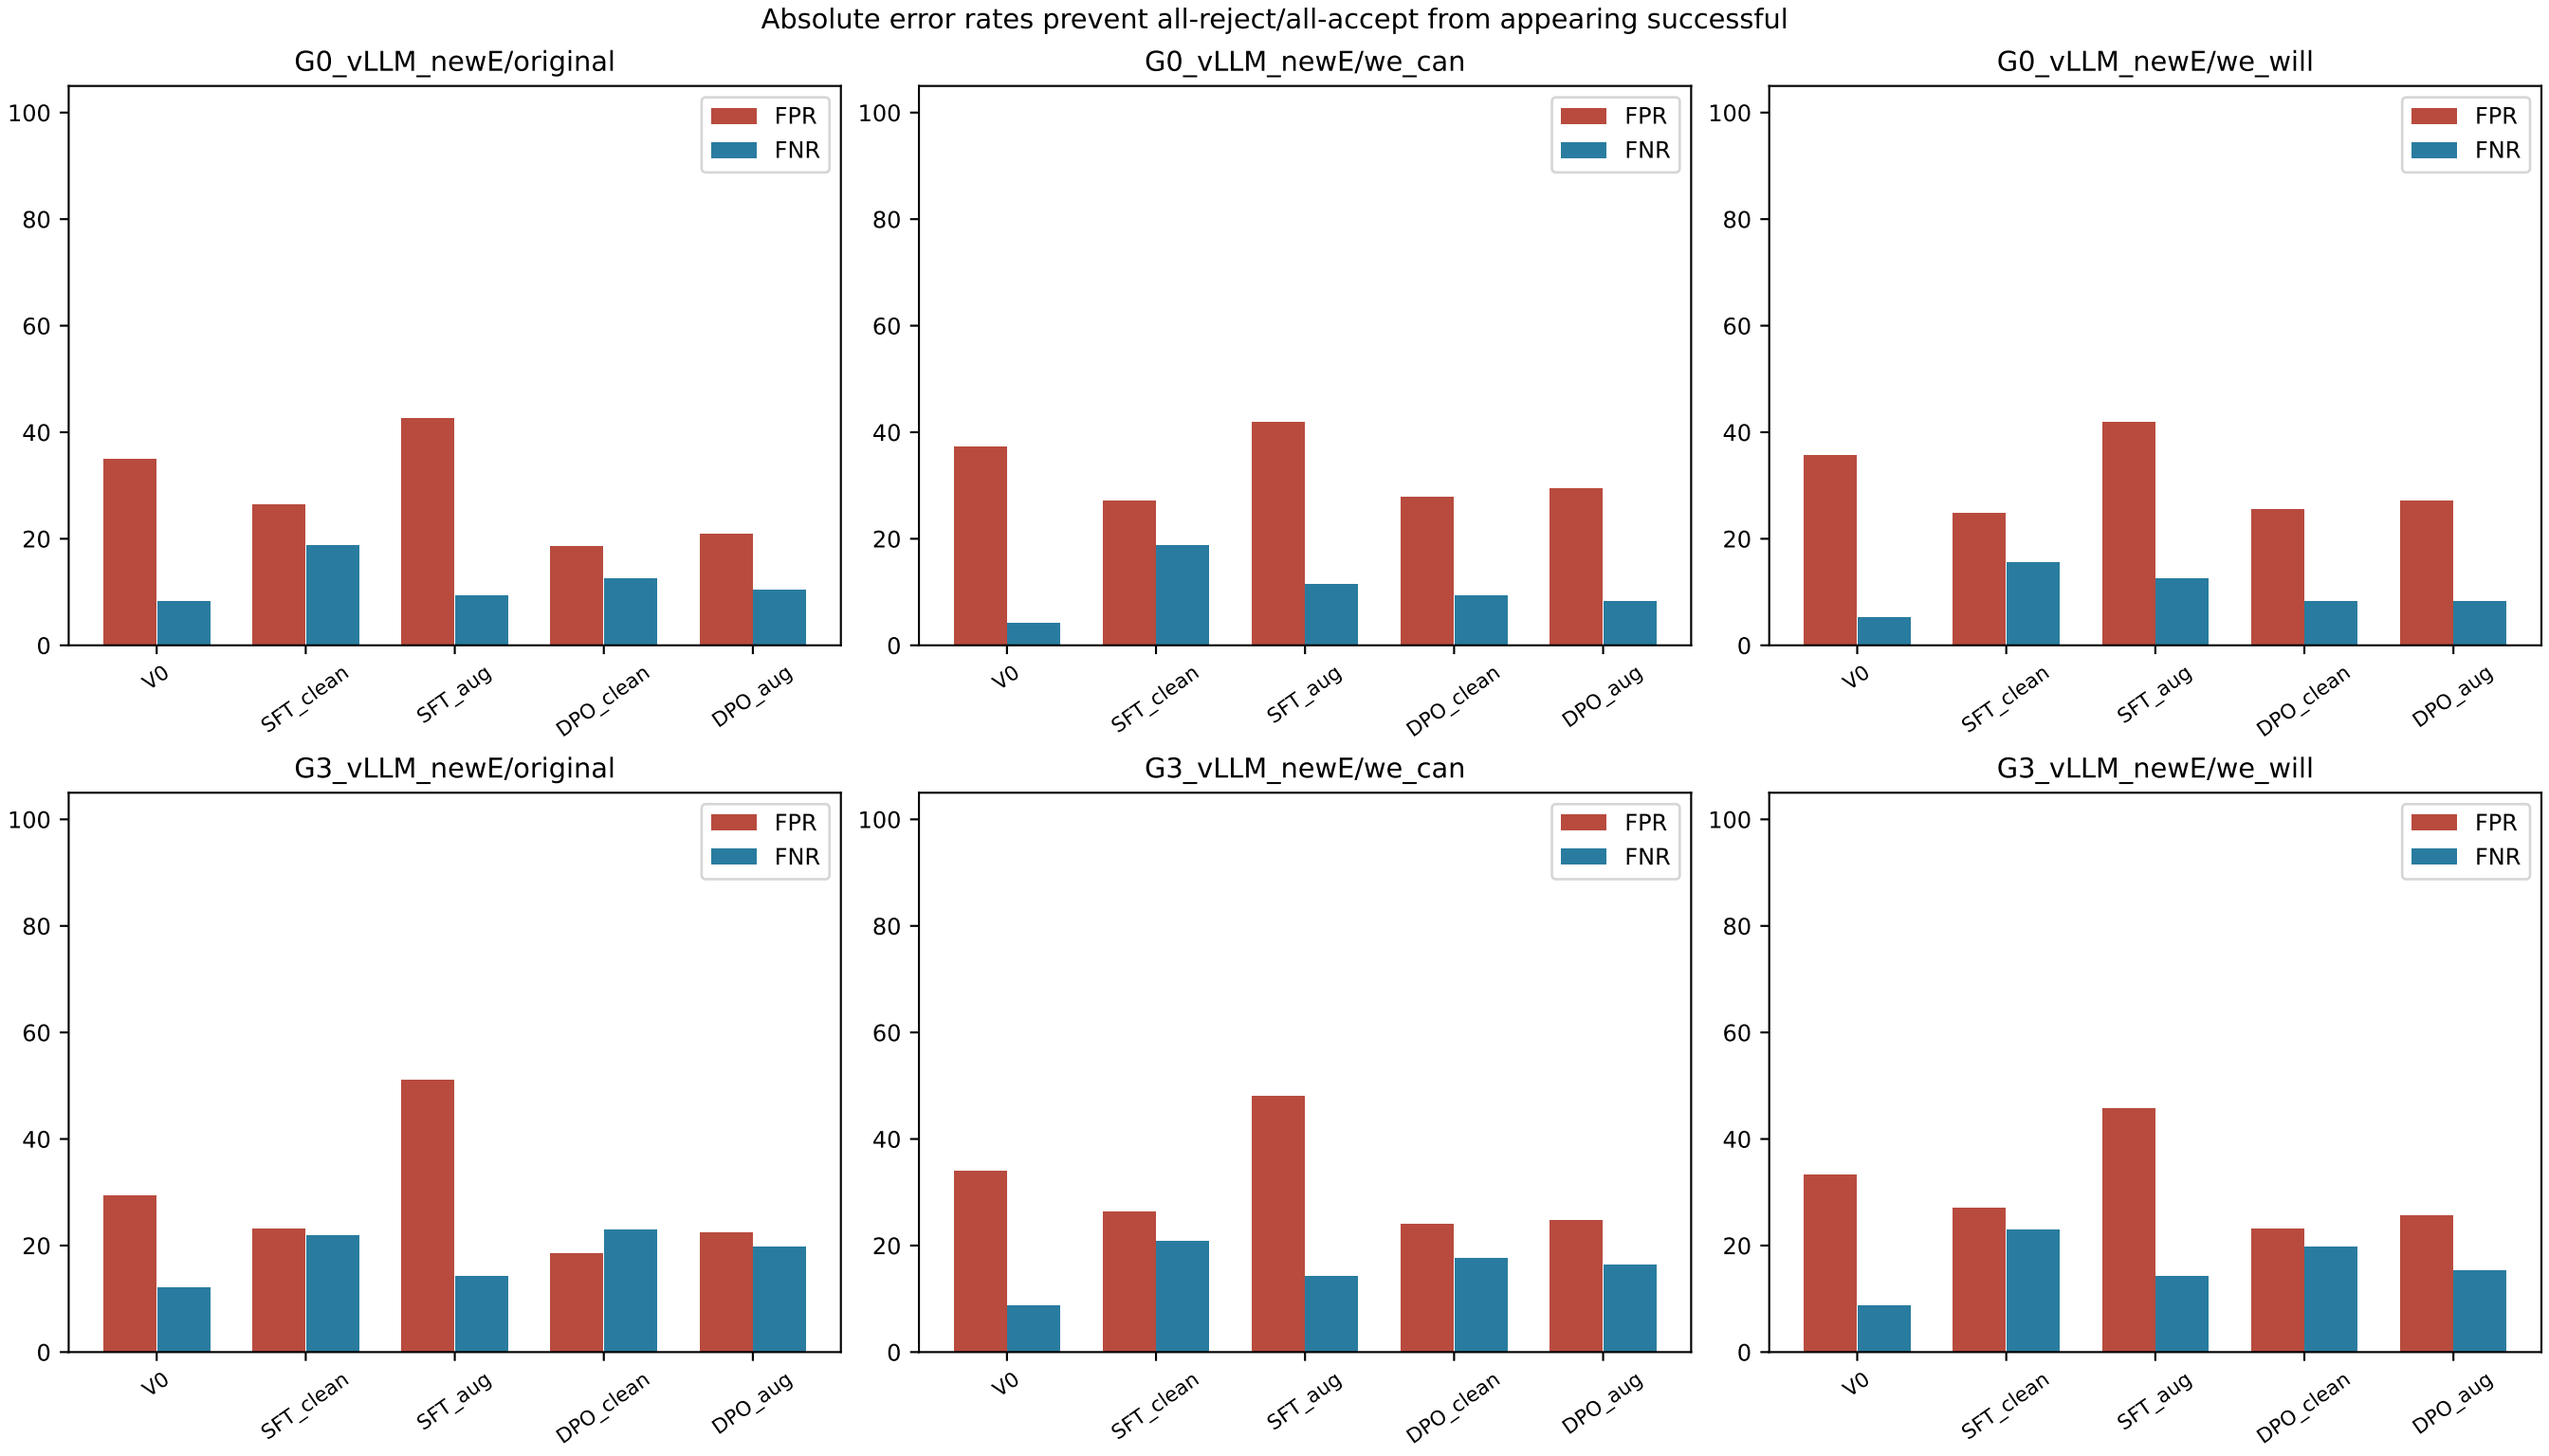
Absolute error rates prevent all-reject/all-accept from appearing successful
G0_vLLM_newE/original
G0_vLLM_newE/we_can
G0_vLLM_newE/we_will
G3_vLLM_newE/original
G3_vLLM_newE/we_can
G3_vLLM_newE/we_will
FPR
FNR
0
20
40
60
80
100
V0
SFT_clean
SFT_aug
DPO_clean
DPO_aug

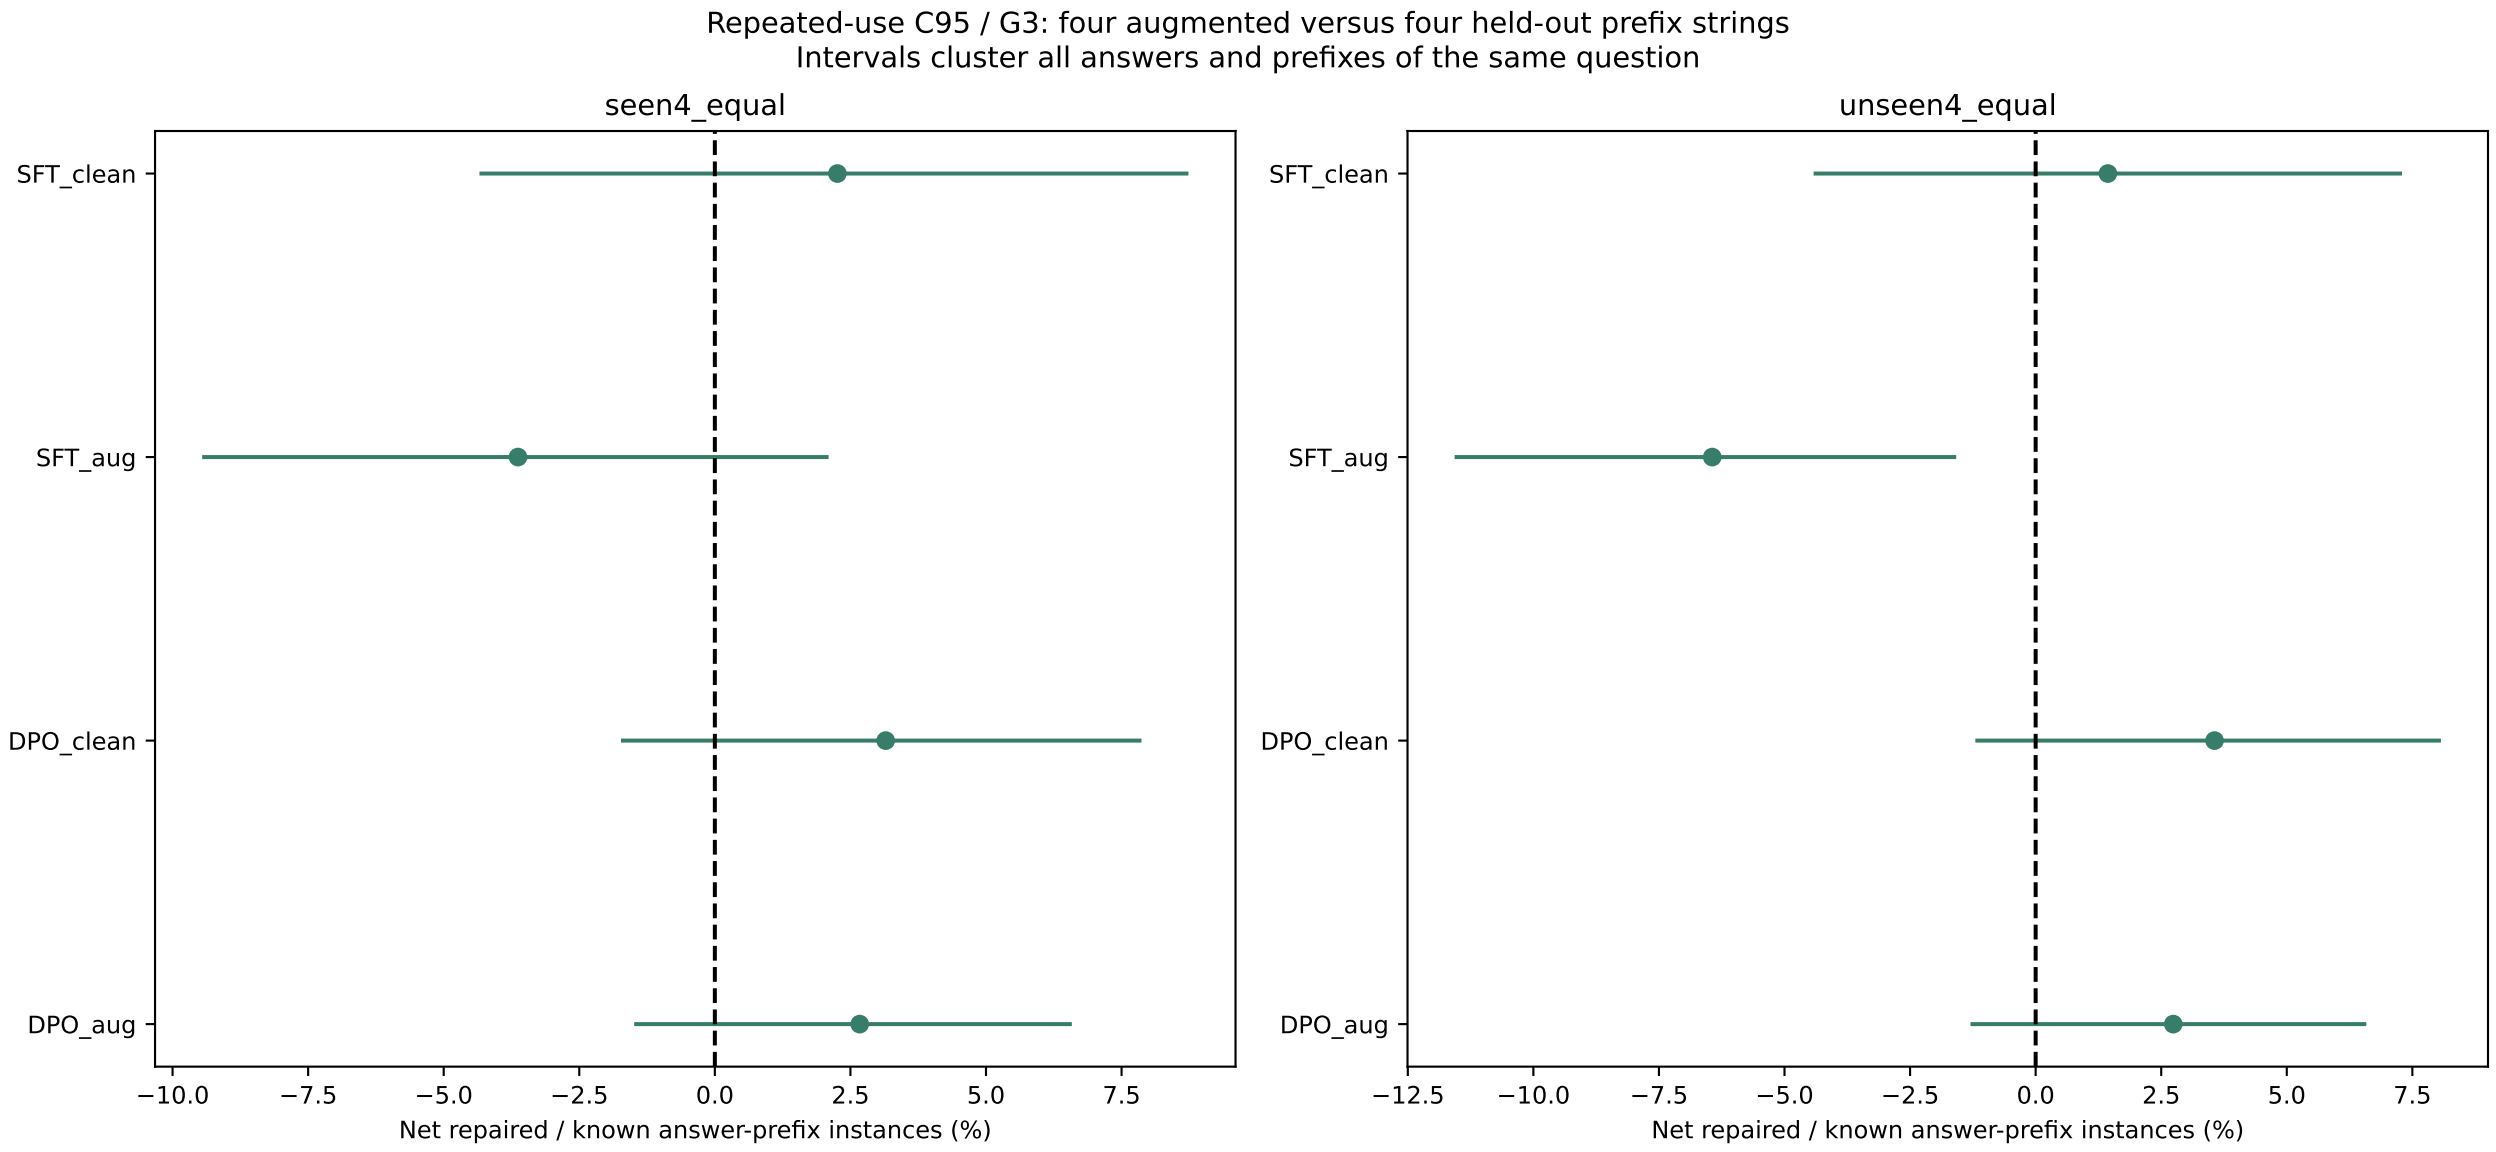
Repeated-use C95 / G3: four augmented versus four held-out prefix strings
Intervals cluster all answers and prefixes of the same question
seen4_equal
unseen4_equal
SFT_clean
SFT_aug
DPO_clean
DPO_aug
−12.5
−10.0
−7.5
−5.0
−2.5
0.0
2.5
5.0
7.5
Net repaired / known answer-prefix instances (%)

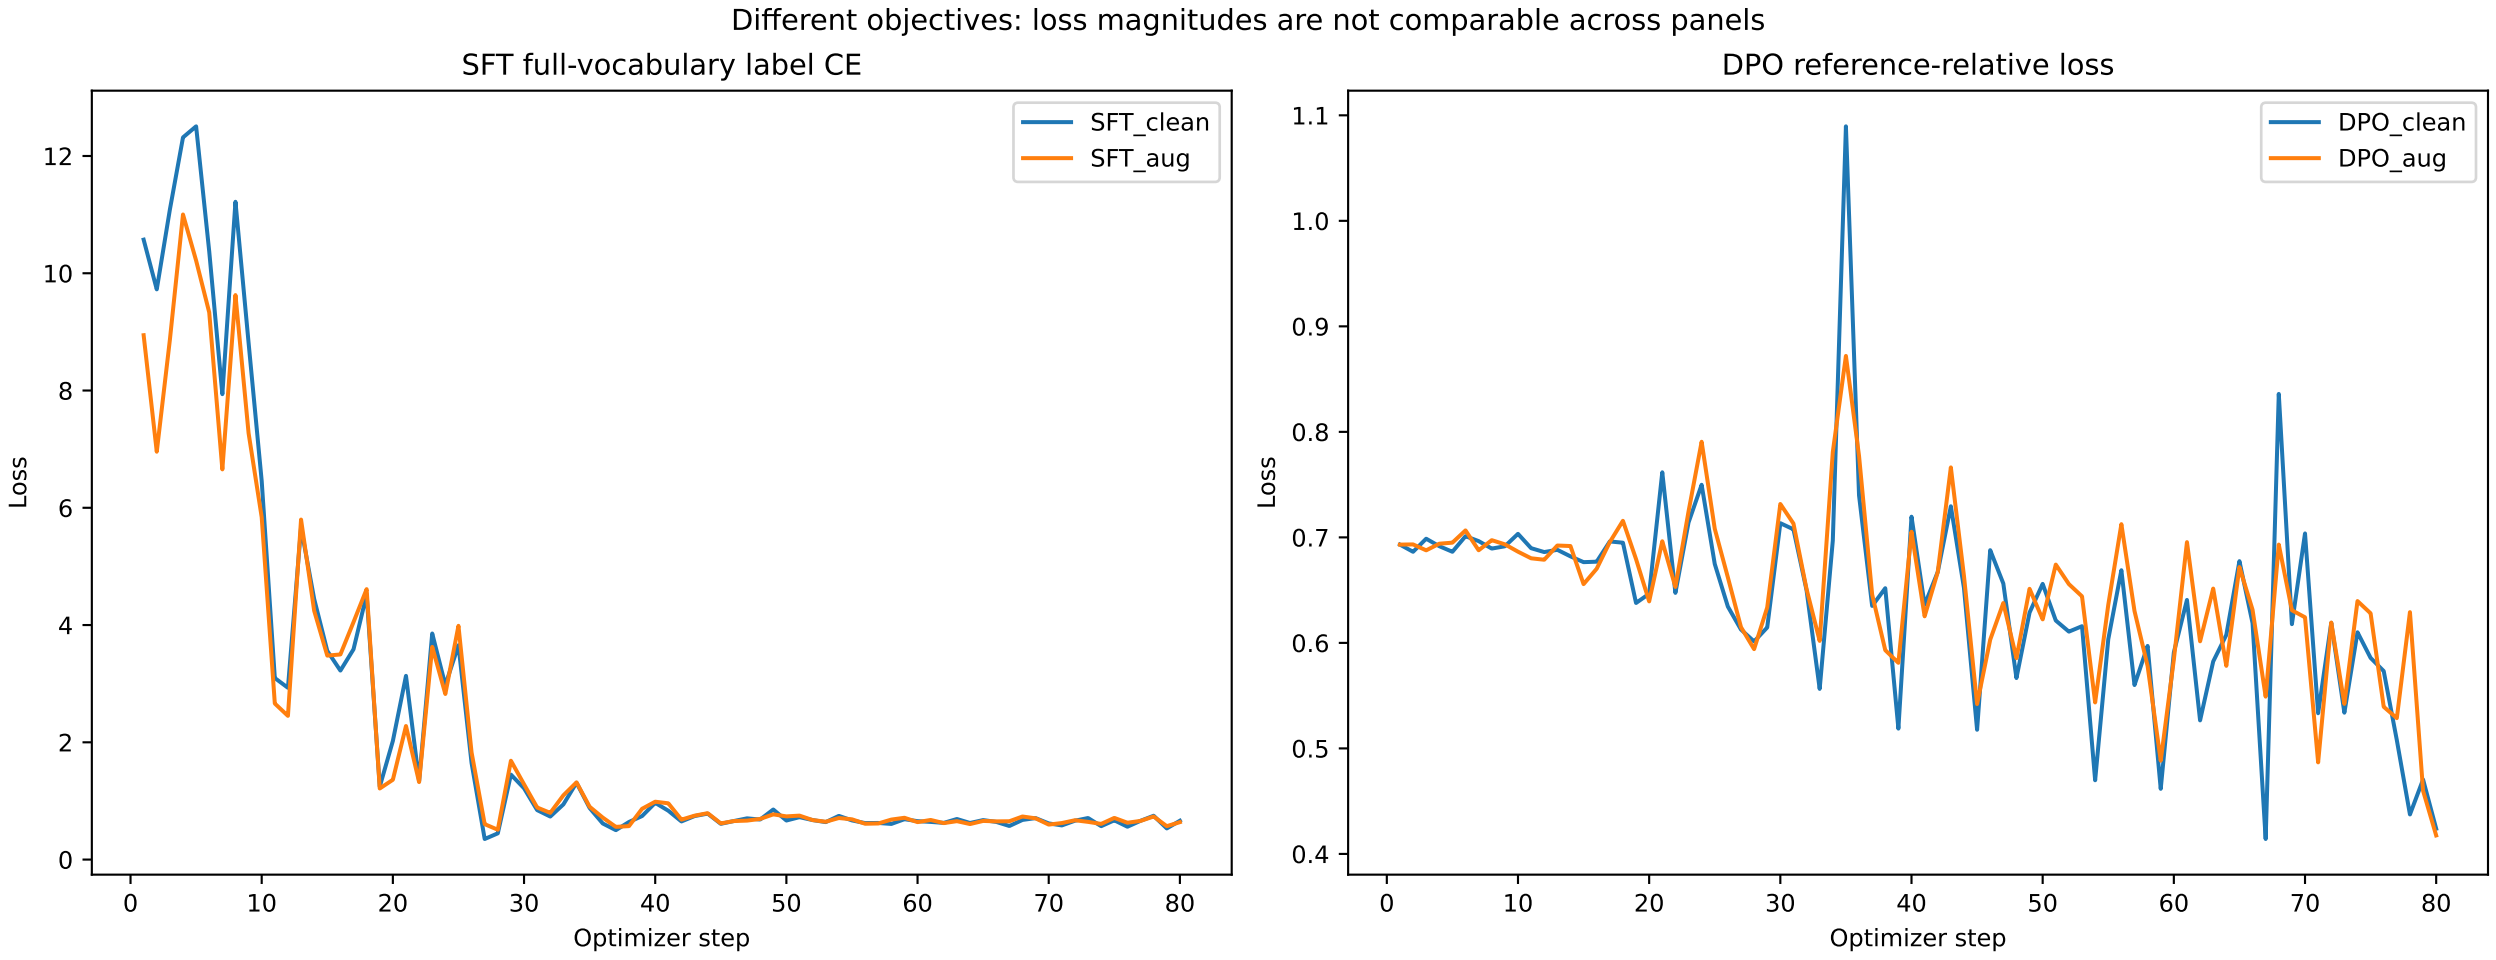
Different objectives: loss magnitudes are not comparable across panels
SFT full-vocabulary label CE
SFT_clean
SFT_aug
Loss
Optimizer step
DPO reference-relative loss
DPO_clean
DPO_aug
Loss
Optimizer step