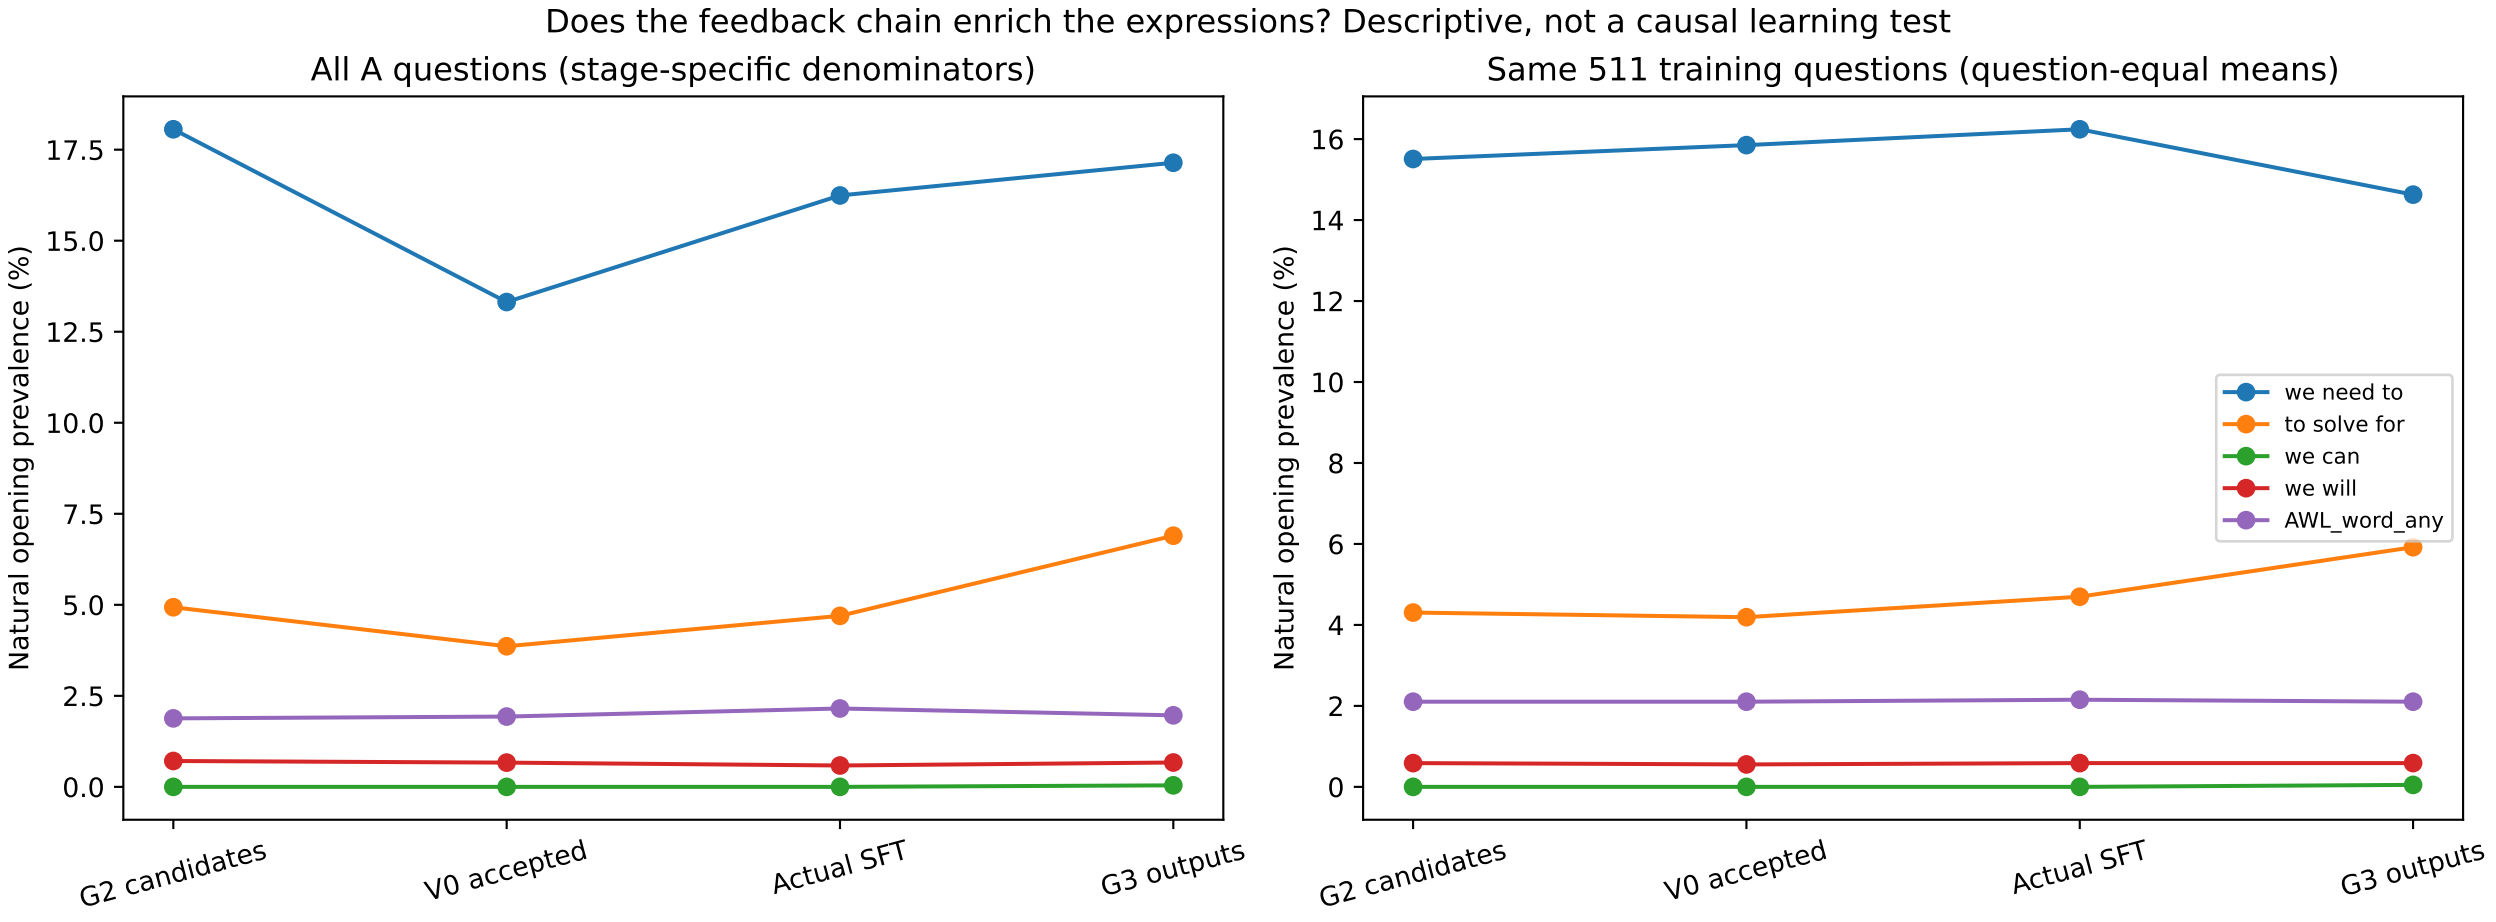
Does the feedback chain enrich the expressions? Descriptive, not a causal learning test
All A questions (stage-specific denominators)
Same 511 training questions (question-equal means)
Natural opening prevalence (%)
Natural opening prevalence (%)
0.0
2.5
5.0
7.5
10.0
12.5
15.0
17.5
0
2
4
6
8
10
12
14
16
G2 candidates
V0 accepted
Actual SFT
G3 outputs
G2 candidates
V0 accepted
Actual SFT
G3 outputs
we need to
to solve for
we can
we will
AWL_word_any

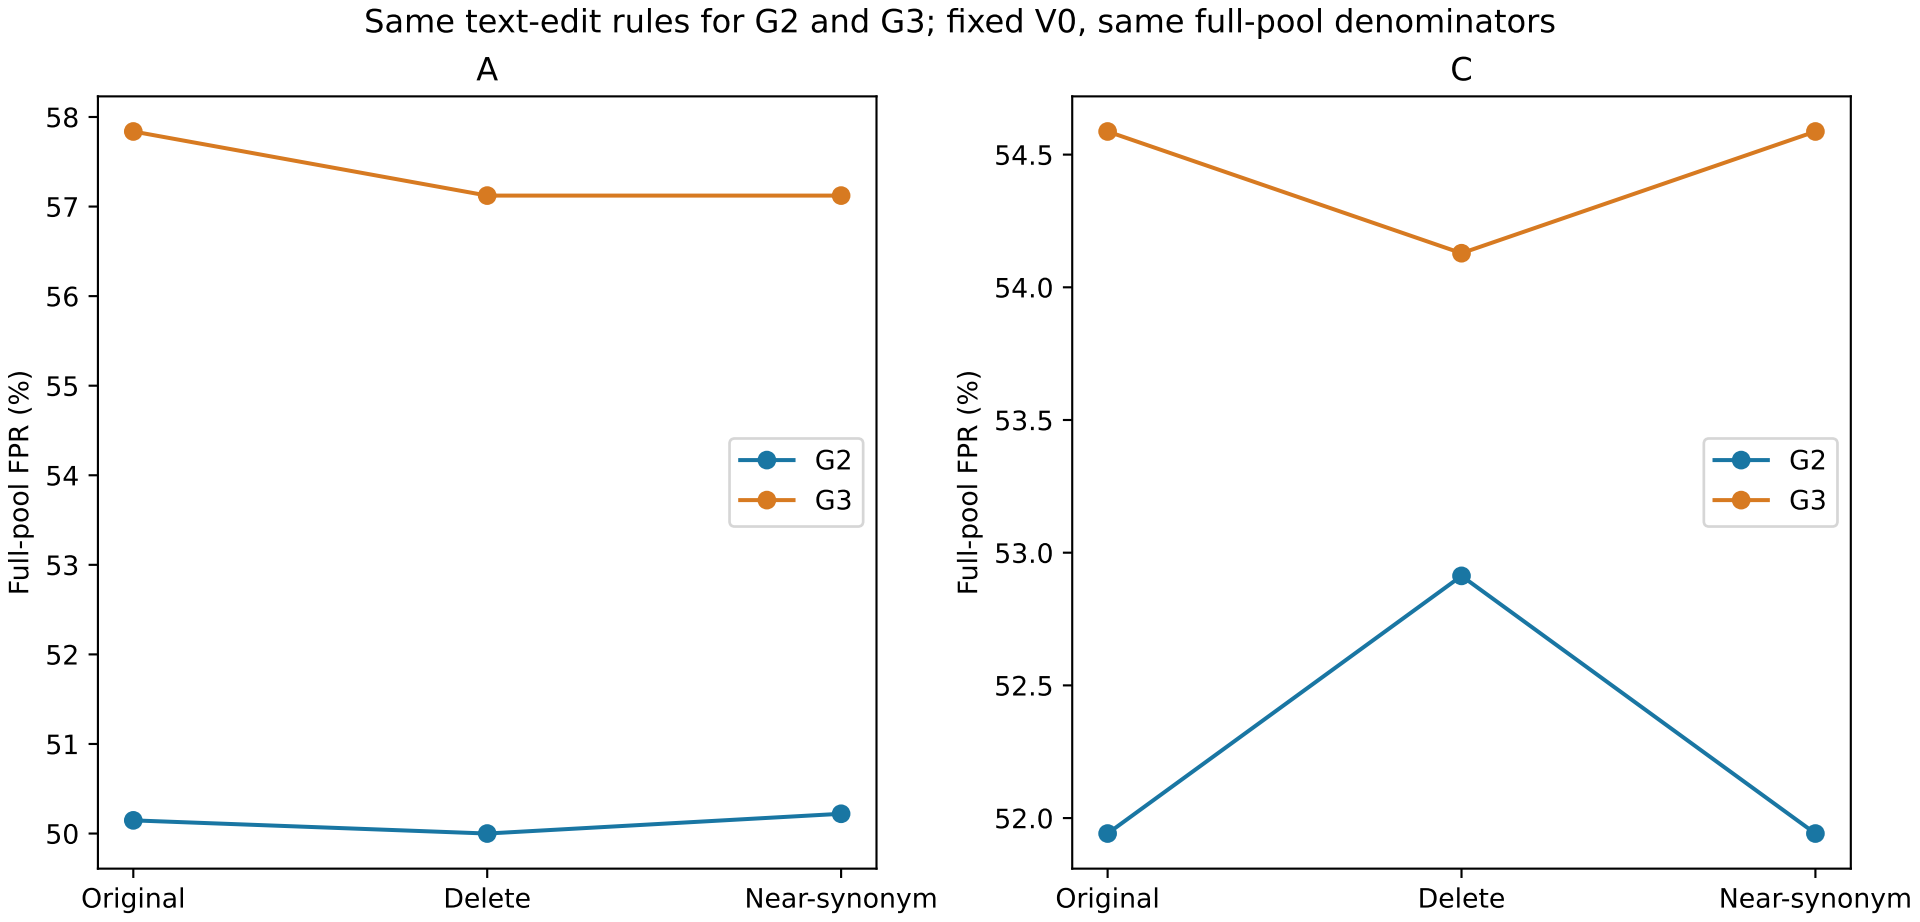
Same text-edit rules for G2 and G3; fixed V0, same full-pool denominators
A
C
Full-pool FPR (%)
58
57
56
55
54
53
52
51
50
54.5
54.0
53.5
53.0
52.5
52.0
G2
G3
Original
Delete
Near-synonym

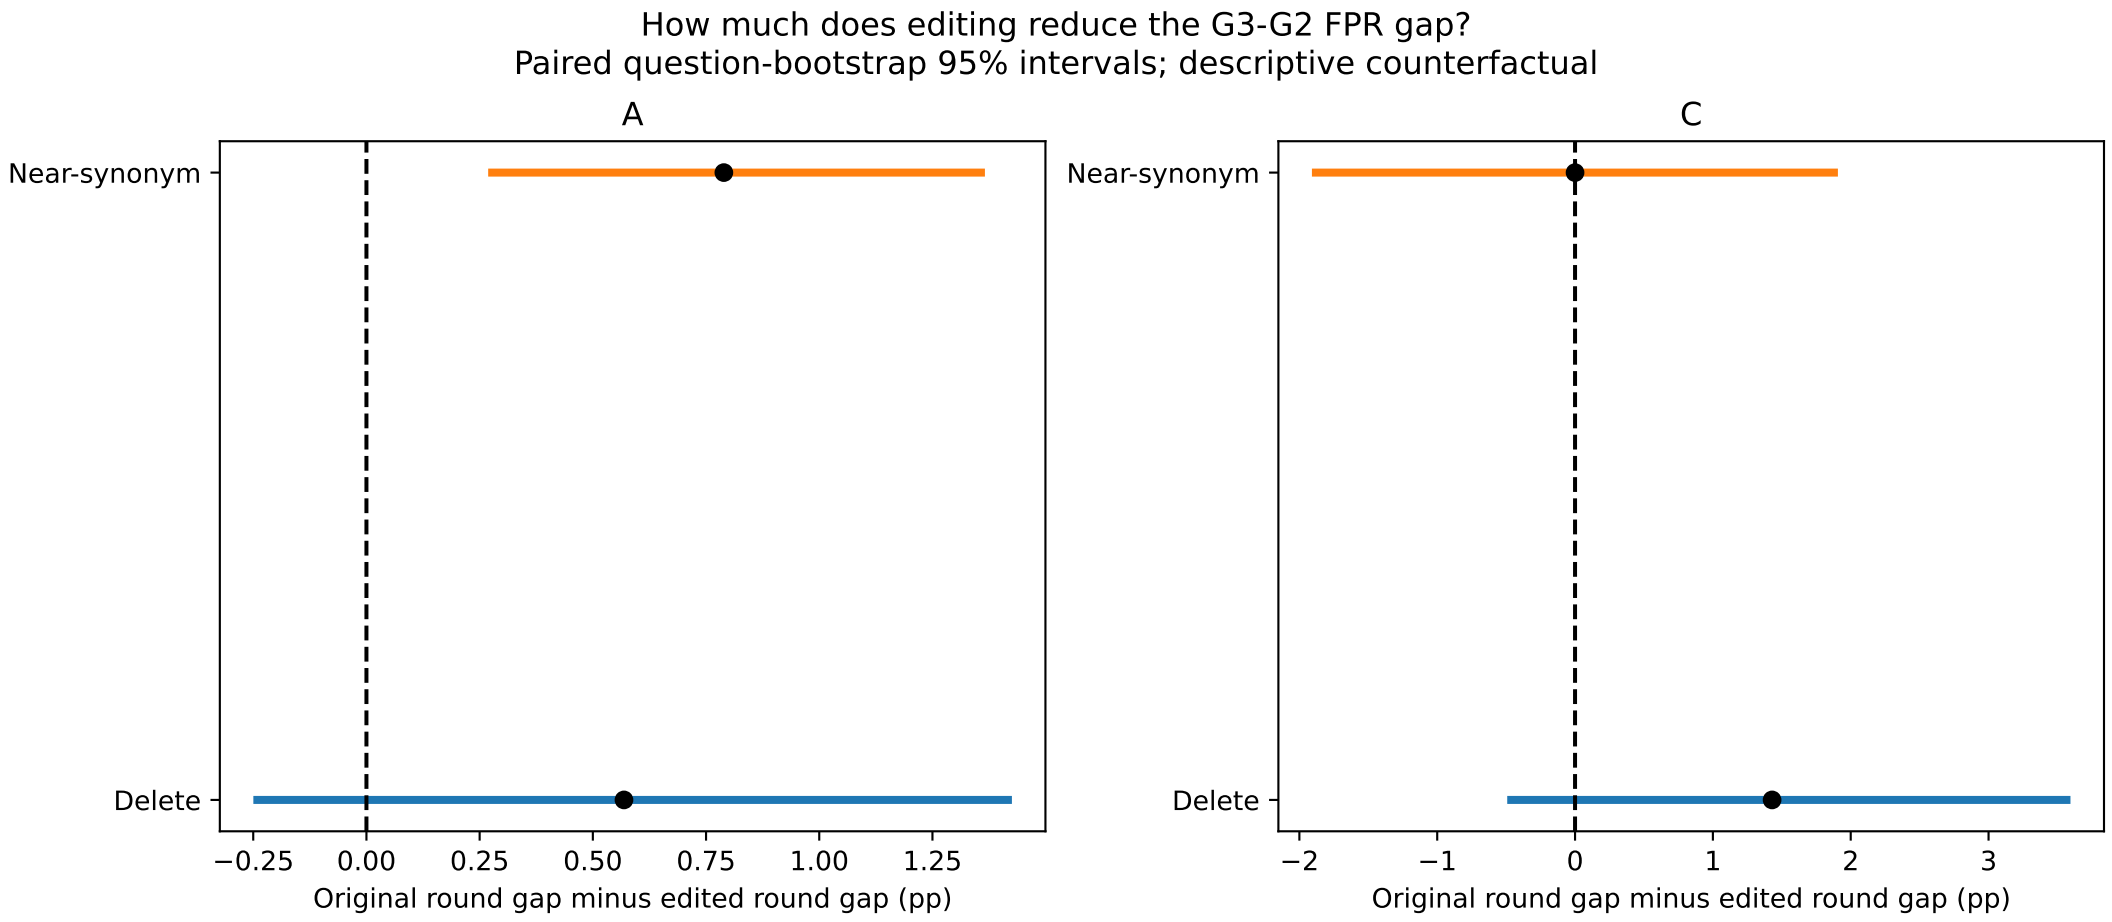
How much does editing reduce the G3-G2 FPR gap?
Paired question-bootstrap 95% intervals; descriptive counterfactual
A
C
Near-synonym
Delete
−0.25
0.00
0.25
0.50
0.75
1.00
1.25
−2
−1
0
1
2
3
Original round gap minus edited round gap (pp)
Original round gap minus edited round gap (pp)

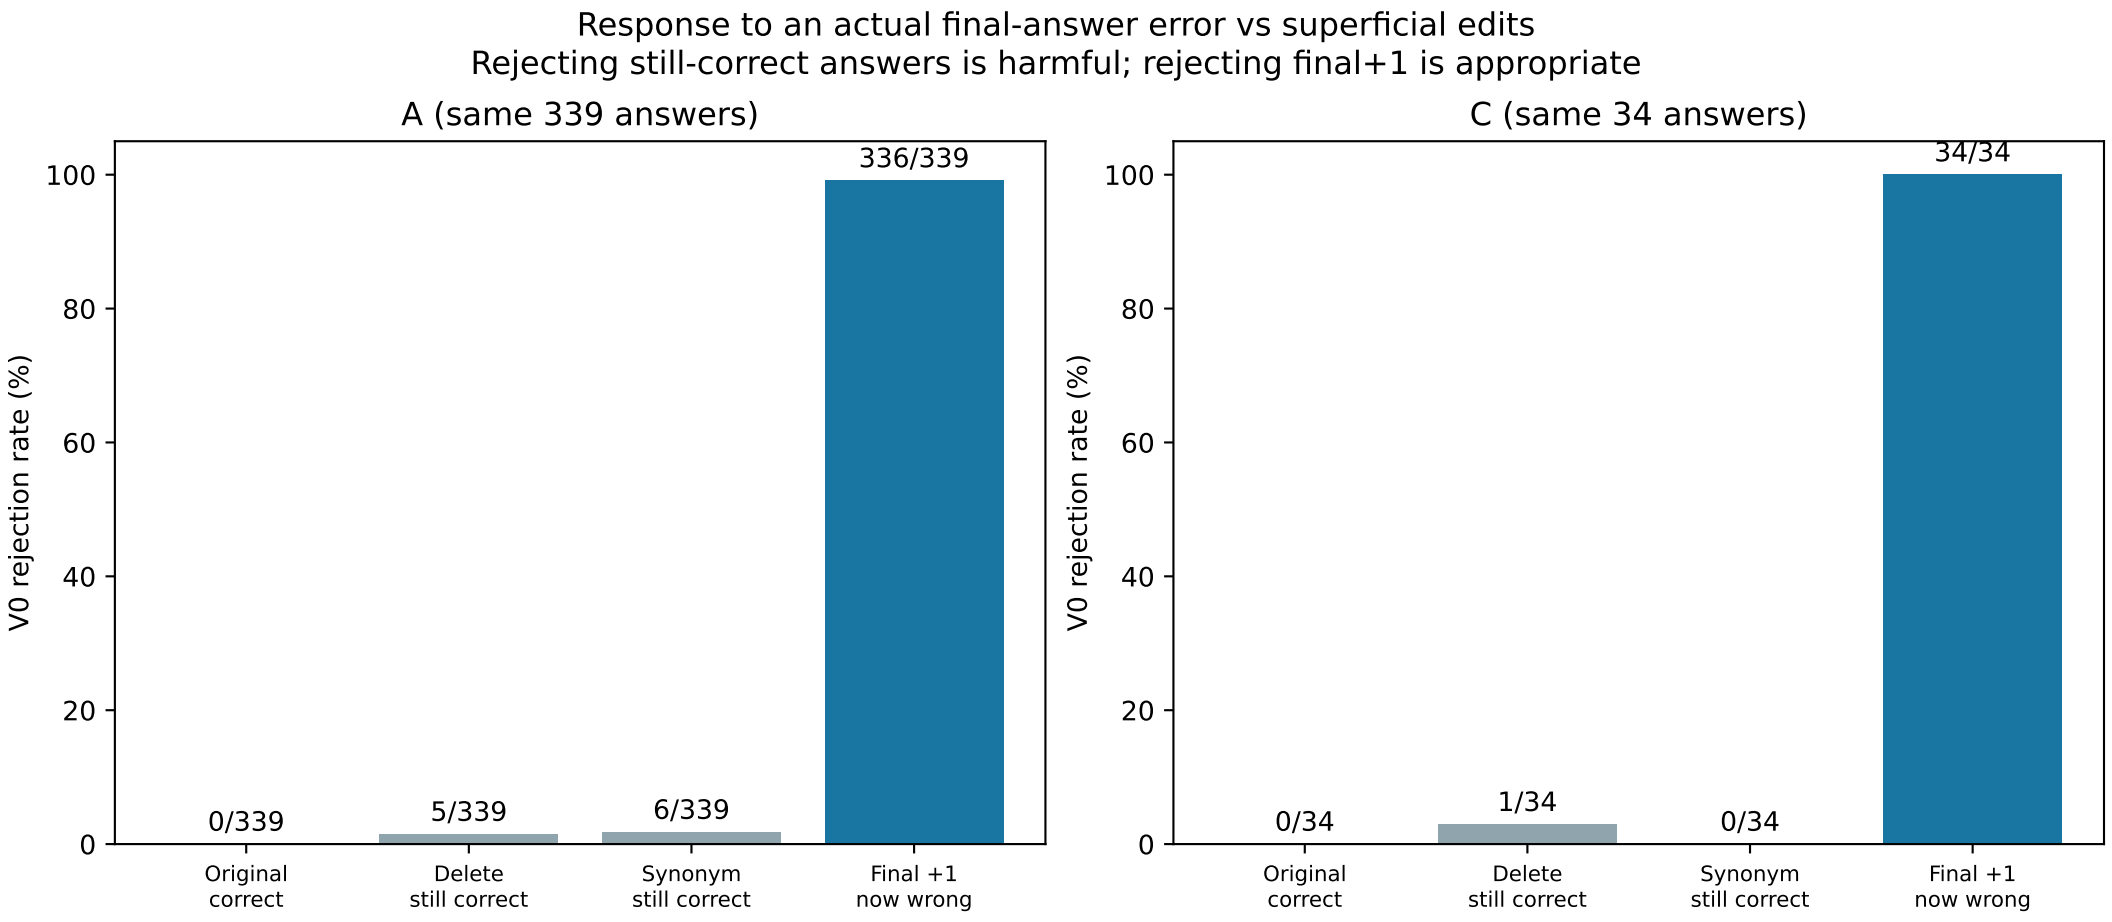
Response to an actual final-answer error vs superficial edits
Rejecting still-correct answers is harmful; rejecting final+1 is appropriate
A (same 339 answers)
C (same 34 answers)
V0 rejection rate (%)
0
20
40
60
80
100
0/339
5/339
6/339
336/339
0/34
1/34
0/34
34/34
Original correct
Delete still correct
Synonym still correct
Final +1 now wrong